Slides and Resources

Post Hoc, Local and Model-Agnostic Explanations

Yilun Zhou

MIT CSAIL & Amazon

Outline

Slides and Resources

- Why model explanations?
- How to compute model explanations? -- Definitions
- How to evaluate model explanations? -- Evaluations
- A definition-evaluation duality

Why Model Explanations?

Slides and Resources

Why Model Explanations?

Slides and Resources

Task for DNN	Caption image	Recognise pneumonia
Problem	Describes green hillside as grazing sheep	Fails on scans from new hospitals
Shortcut	Uses background to recognise primary object	Looks at hospital token, not lung

(Geirhos et al. 2020)

Why Model Explanations?

Slides and Resources

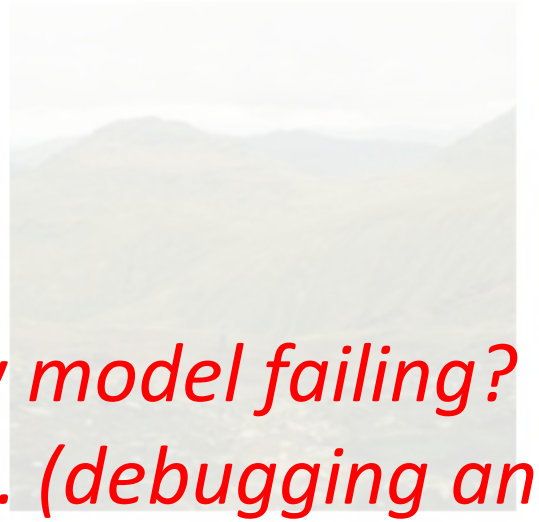

Why is my model failing?
Because … (debugging and diagnosis)

Task for DNN	Caption image	Recognise pneumonia
Problem	Describes green hillside as grazing sheep	Fails on scans from new hospitals
Shortcut	Uses background to recognise primary object	Looks at hospital token, not lung

(Geirhos et al. 2020)

Why Model Explanations?

Slides and Resources

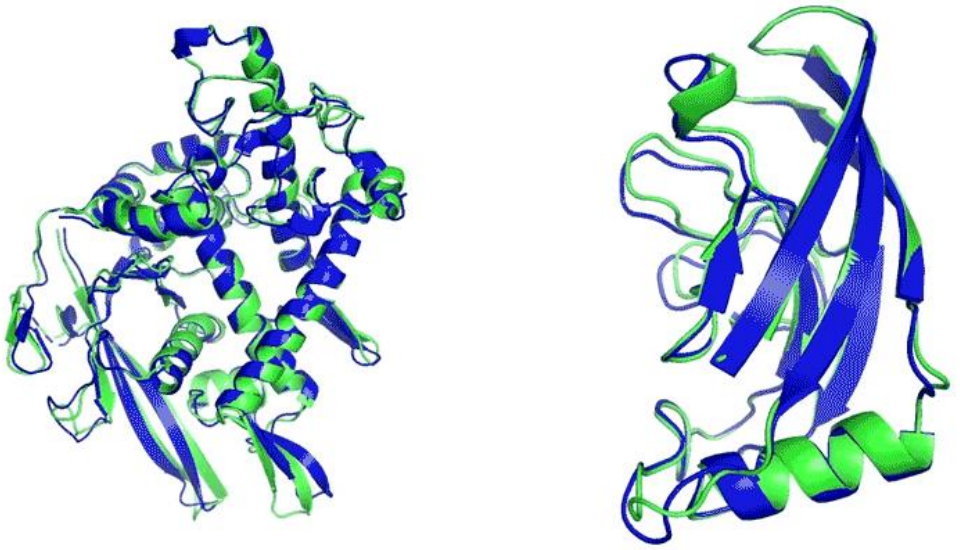

Why Model Explanations?

Slides and Resources

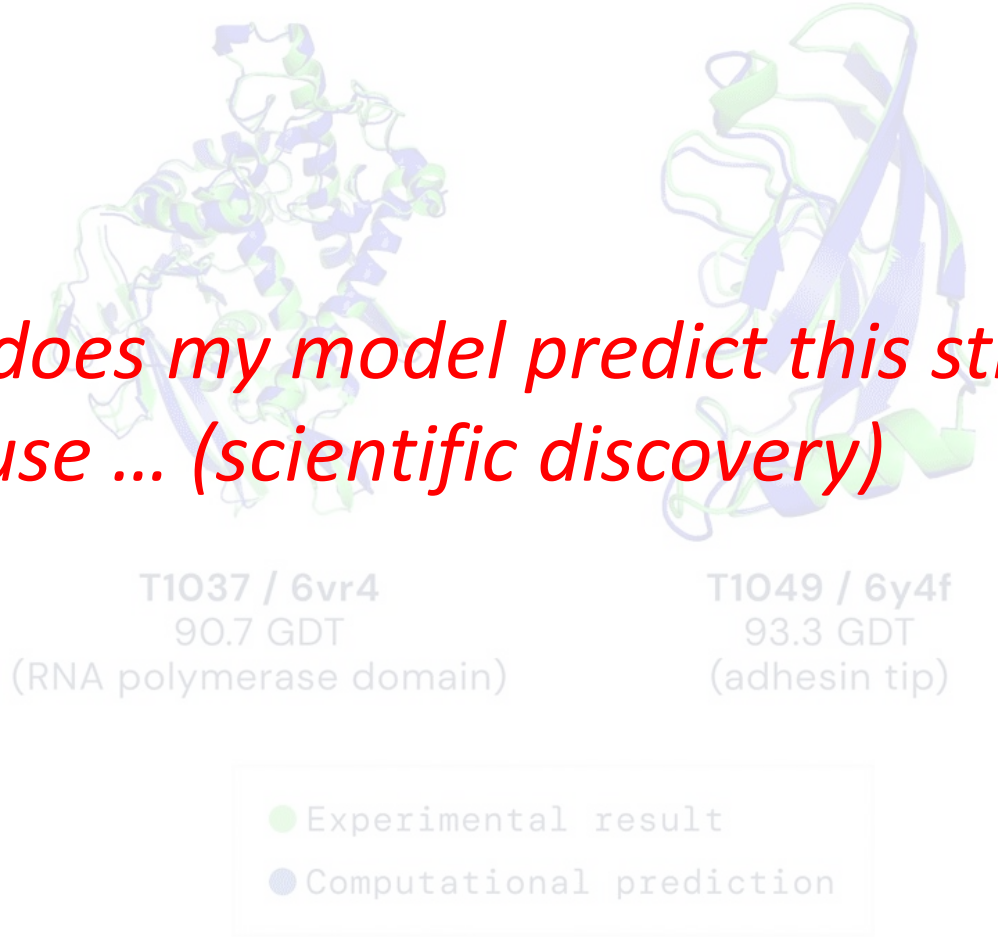

How does my model predict this structure?
Because … (scientific discovery)

The Two Major Axes of Interpretability

Slides and Resources

Global
Explains the model “more generally”

Local
Requires an input instance

The Two Major Axes of Interpretability

Slides and Resources

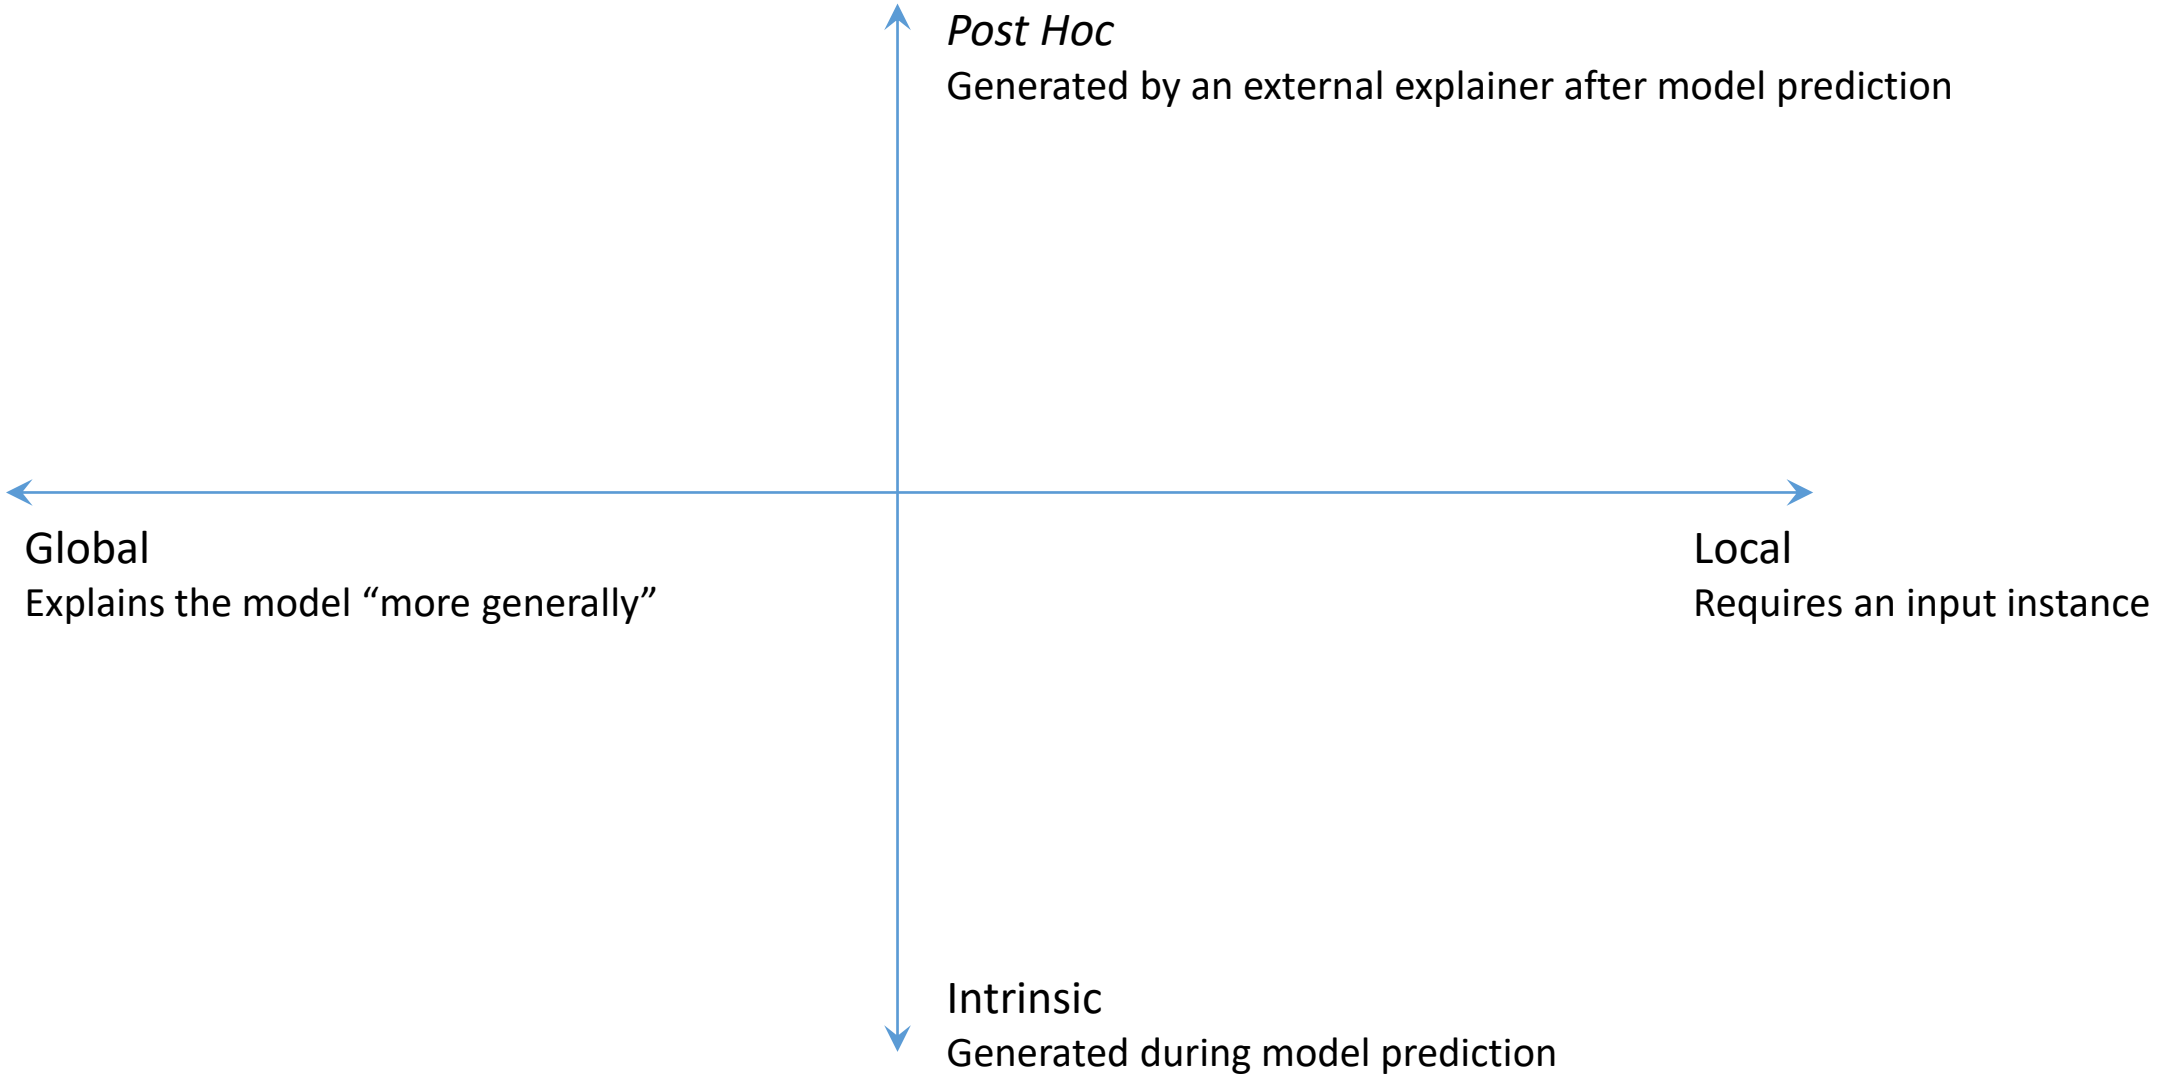

The Two Major Axes of Interpretability

Slides and Resources

Post Hoc
Generated by an external explainer after model prediction

Global
Explains the model “more generally”

Local
Requires an input instance

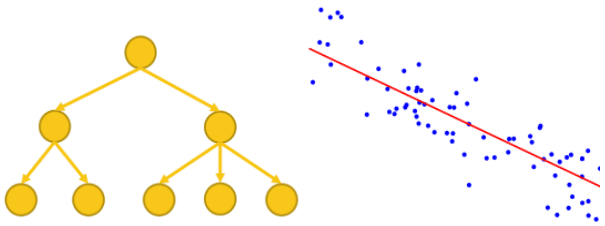

Intrinsic
Generated during model prediction

The Two Major Axes of Interpretability

Slides and Resources

Post Hoc
Generated by an external explainer after model prediction

Global
Explains the model “more generally”

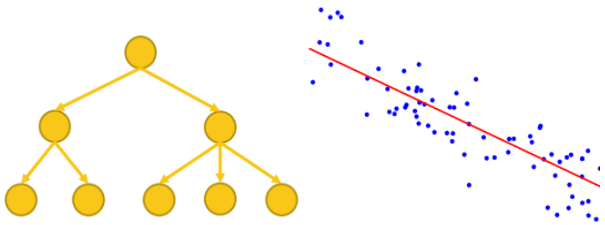

Local
Requires an input instance

Next presentation

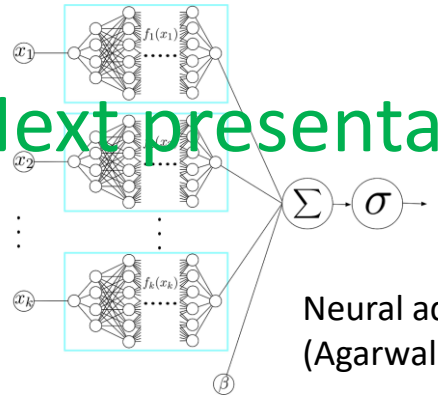

Neural additive model
(Agarwal et al., 2021)

Intrinsic
Generated during model prediction

The Two Major Axes of Interpretability

Slides and Resources

CNN filter visualization
(Olah et al., 2017)

Post Hoc
Generated by an external explainer after model prediction

Global
Explains the model "more generally"

Local
Requires an input instance

Next presentation

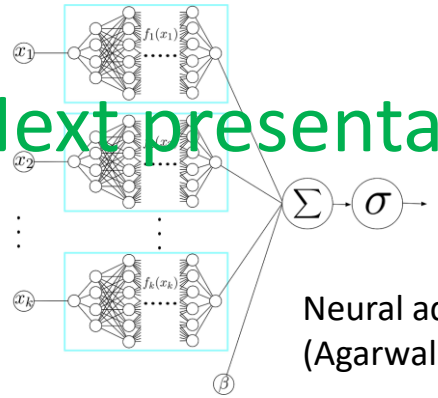

Neural additive model
(Agarwal et al., 2021)

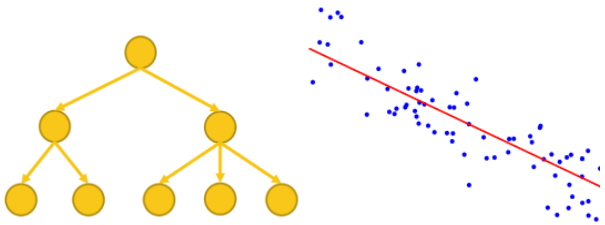

Intrinsic
Generated during model prediction

The Two Major Axes of Interpretability

Slides and Resources

CNN filter visualization
(Olah et al., 2017)

Post Hoc
Generated by an external explainer after model prediction

This presentation

Global
Explains the model "more generally"

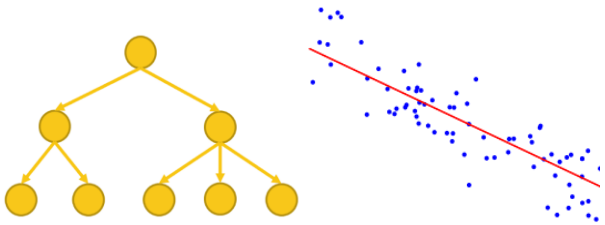

Local
Requires an input instance

Next presentation

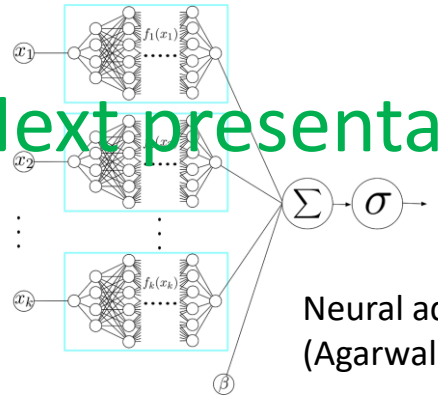

Neural additive model
(Agarwal et al., 2021)

Intrinsic
Generated during model prediction

Slides and Resources

Local *Post Hoc* Explanations

- What is a local *post hoc* explanation?
  - Some description of the model's local decision making logic
- If this image patch is grayed out, the model prediction will change to this.

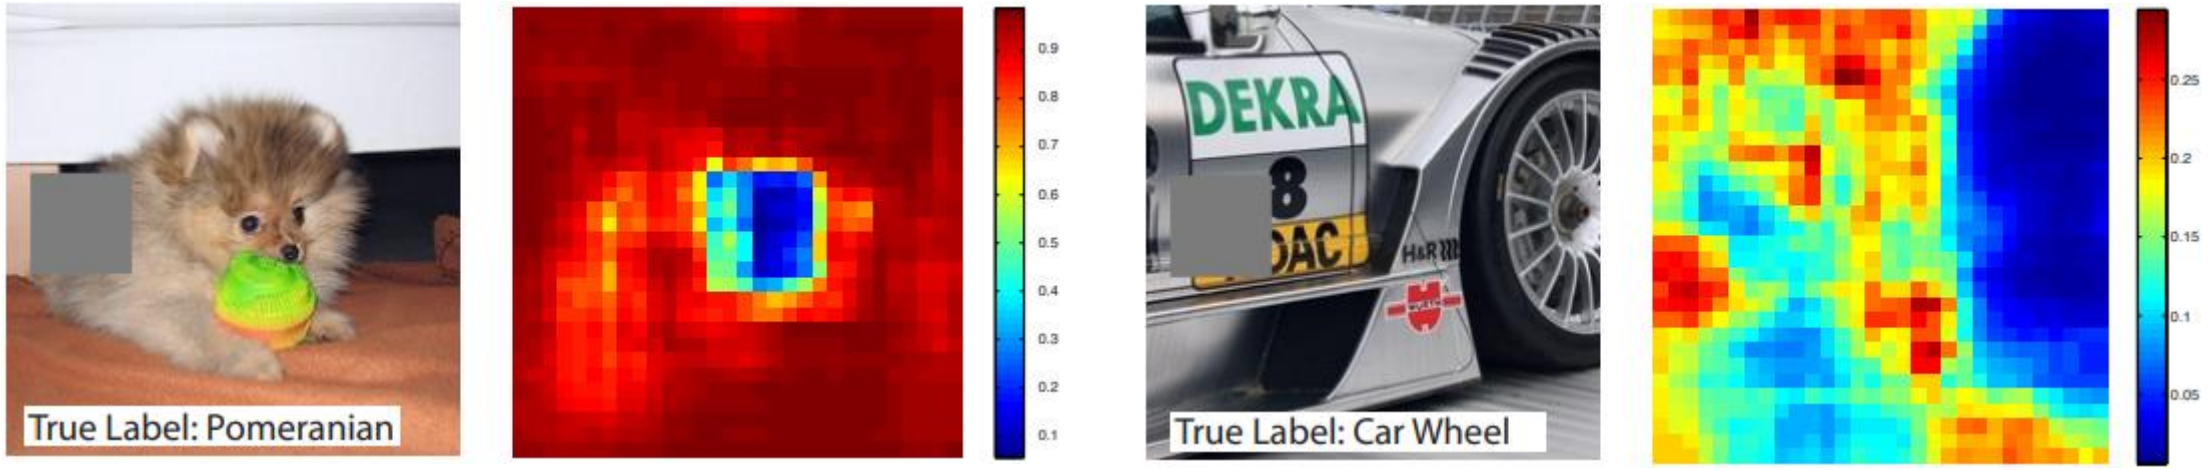

Occlusion Saliency (Zeiler and Fergus, 2014)

Local *Post Hoc* Explanations

Slides and Resources

- What is a local *post hoc* explanation?
  - Some description of the model's local decision making logic
- If this input feature is changed to this value, the model prediction will be different.

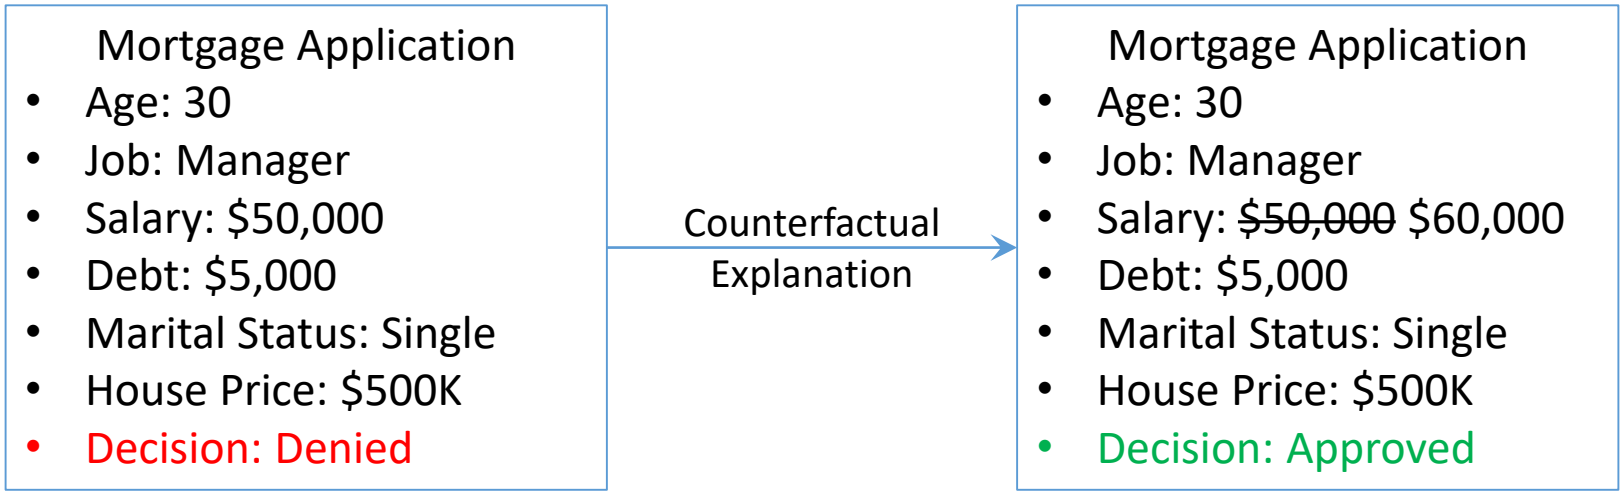

Local *Post Hoc* Explanations

Slides and Resources

- What is a local *post hoc* explanation?
  - Some description of the model's local decision making logic
- If this training instance is not present in the dataset, the model will make a different prediction.

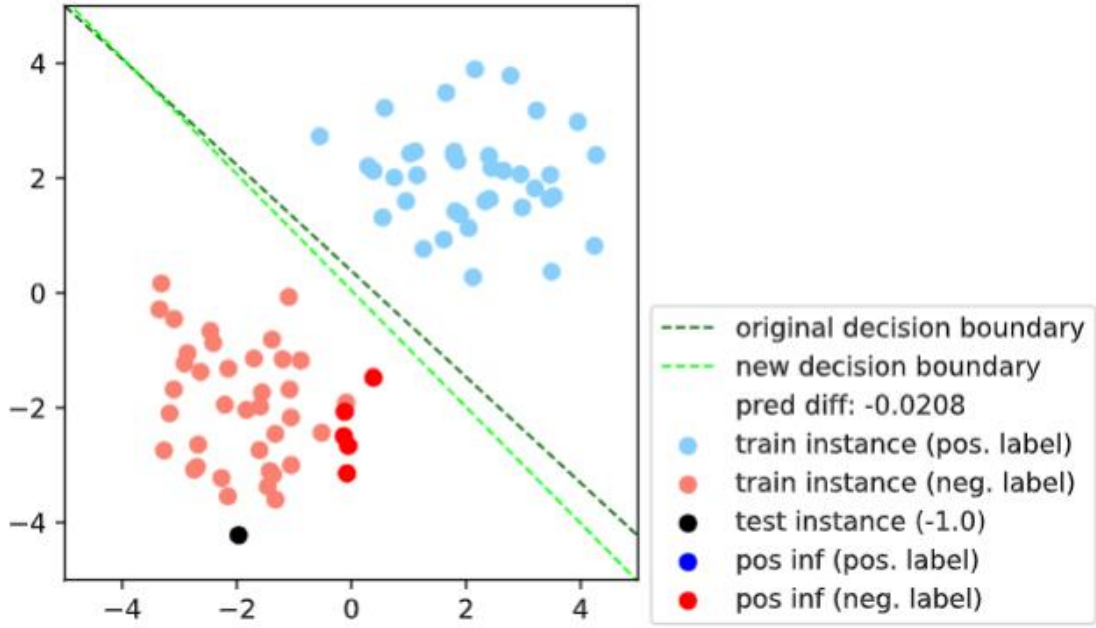

Feature Attributions (Saliency Maps)

Slides and Resources

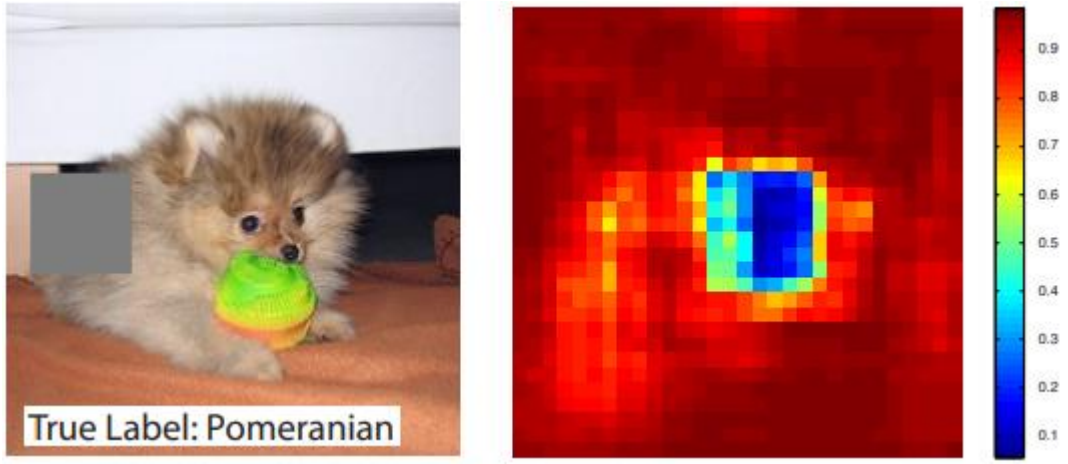

Occlusion Saliency (Zeiler and Fergus, 2014)

Feature Attributions (Saliency Maps)

Slides and Resources

- Vanilla gradient: $s = \nabla_x f(x)$

Feature Attributions (Saliency Maps)

Slides and Resources

- Vanilla gradient: $s = \nabla_x f(x)$
- SmoothGrad: $s = \mathbb{E}_\epsilon[\nabla_x f(x + \epsilon)]$; ϵ_i independent for every x_i

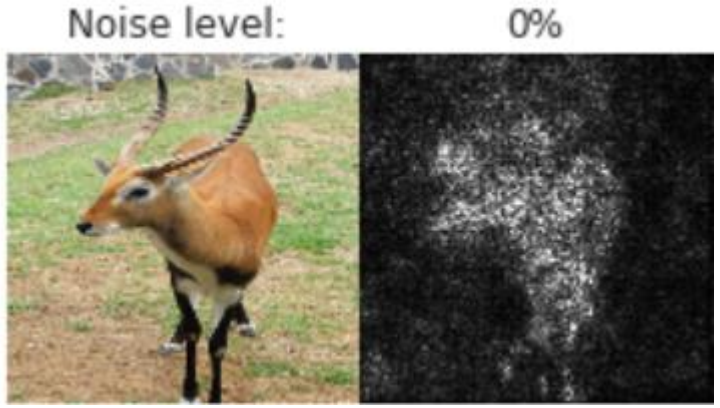

Feature Attributions (Saliency Maps)

Slides and Resources

- Vanilla gradient: $s = \nabla_x f(x)$
- SmoothGrad: $s = \mathbb{E}_\epsilon[\nabla_x f(x + \epsilon)]$; ϵ_i independent for every x_i

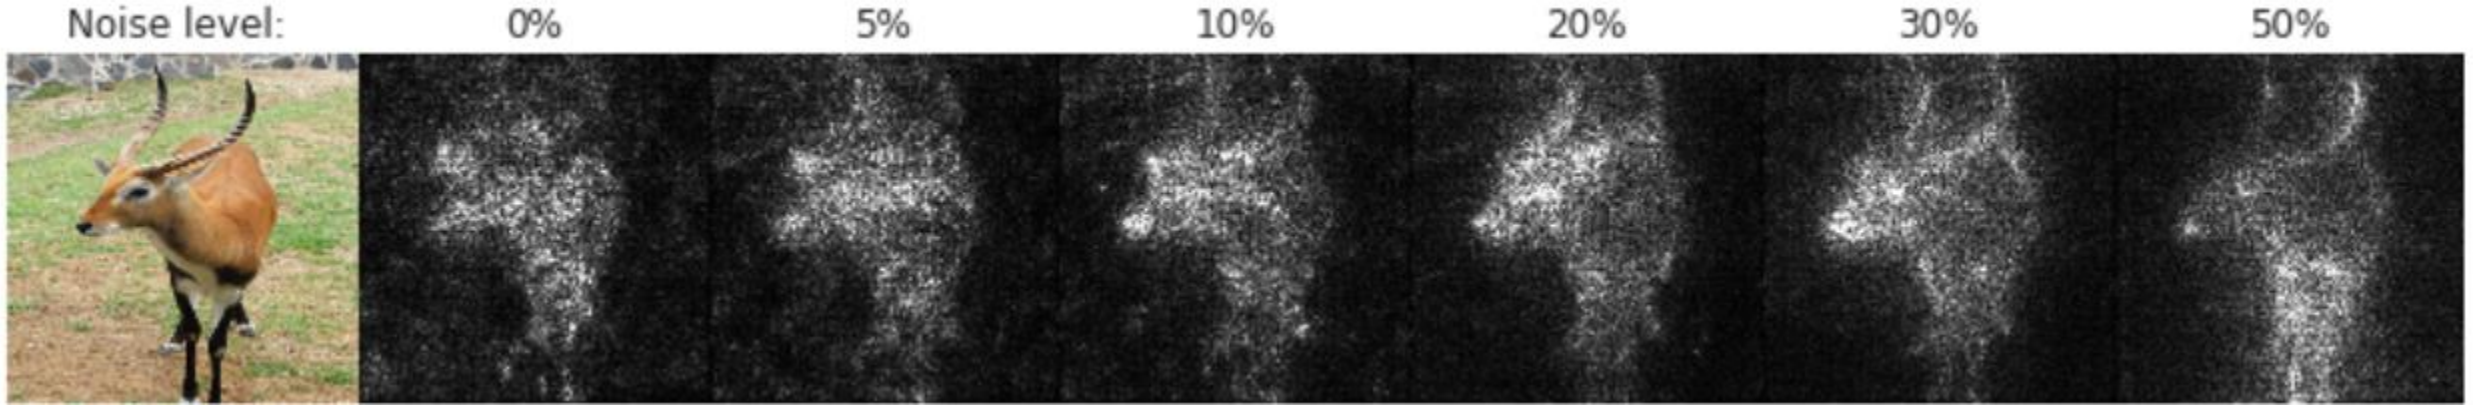

Feature Attributions (Saliency Maps)

Slides and Resources

- Vanilla gradient: $s = \nabla_x f(x)$
- SmoothGrad: $s = \mathbb{E}_\epsilon[\nabla_x f(x+\epsilon)]$; ϵ_i independent for every x_i
- Integrated gradient: $s = \int_0^1 \nabla_x f\big(x_0 + u(x - x_0)\big)\,\mathrm{d}u$

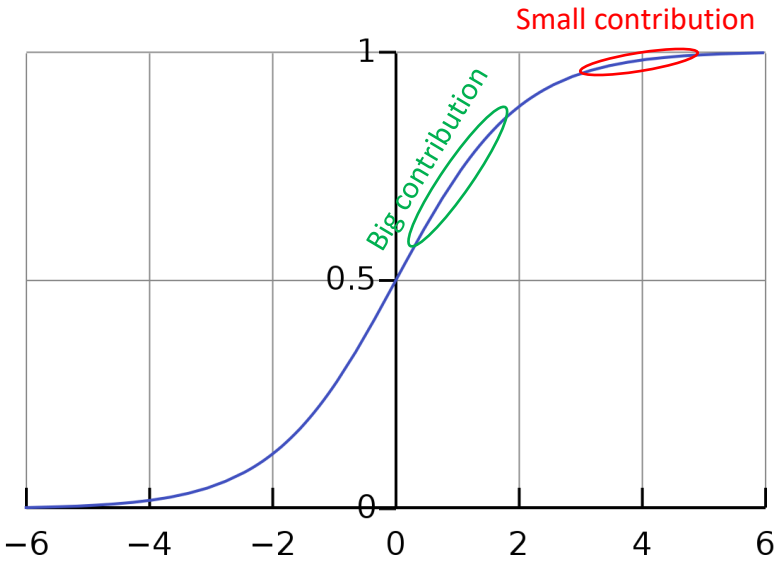

Feature Attributions (Saliency Maps)

Slides and Resources

- Vanilla gradient: $s = \nabla_x f(x)$
- SmoothGrad: $s = \mathbb{E}_\epsilon[\nabla_x f(x+\epsilon)]$; ϵ_i independent for every x_i
- Integrated gradient: $s = \int_0^1 \nabla_x f\big(x_0 + u(x - x_0)\big)\,\mathrm{d}u$
- Occlusion: $s_i = f(x) - f(x_{-i})$

Feature Attributions (Saliency Maps)

Slides and Resources

- Vanilla gradient: $s = \nabla_x f(x)$
- SmoothGrad: $s = \mathbb{E}_\epsilon[\nabla_x f(x+\epsilon)]$; ϵ_i independent for every x_i
- Integrated gradient: $s = \int_0^1 \nabla_x f\big(x_0 + u(x - x_0)\big)\,\mathrm{d}u$
- Occlusion: $s_i = f(x) - f(x_{-i})$
- LIME: $f \sim s^T m + b$

Feature mask m	Masked sentence	$f(\cdot)$
0, 0, 0, 0	""	0.49
0, 0, 0, 1	"beautiful"	0.9
0, 0, 1, 0	"and"	0.52
0, 0, 1, 1	"and beautiful"	0.91
…		
1, 1, 1, 0	"It's good and"	0.89
1, 1, 1, 1	"It's good and beautiful"	0.95

Feature Attributions (Saliency Maps)

Slides and Resources

- Vanilla gradient: $s = \nabla_x f(x)$
- SmoothGrad: $s = \mathbb{E}_\epsilon[\nabla_x f(x+\epsilon)]$; ϵ_i independent for every x_i
- Integrated gradient: $s = \int_0^1 \nabla_x f\big(x_0 + u(x - x_0)\big)\,\mathrm{d}u$
- Occlusion: $s_i = f(x) - f(x_{-i})$
- LIME: $f \sim s^T m + b$
- SHAP: $s_i = \sum_{S \subseteq F \setminus \{i\}} \frac{(|S|!(|F|-|S|-1)!)}{|F|!} \left[f_{S\cup\{i\}}\big(x_{S\cup\{i\}}\big) - f_S(x_S) \right]$

Feature Attributions (Saliency Maps)

Slides and Resources

- Vanilla gradient: $s = \nabla_x f(x)$
- SmoothGrad: $s = \mathbb{E}_\epsilon[\nabla_x f(x+\epsilon)]$; ϵ_i independent for every x_i
- Integrated gradient: $s = \int_0^1 \nabla_x f\big(x_0 + u(x - x_0)\big)\,\mathrm{d}u$
- Occlusion: $s_i = f(x) - f(x_{-i})$
- LIME: $f \sim s^T m + b$
- SHAP: $s_i = \sum_{S \subseteq F \setminus \{i\}} \frac{(|S|!(|F|-|S|-1)!)}{|F|!}\left[f_{S\cup\{i\}}\big(x_{S\cup\{i\}}\big) - f_S(x_S)\right]$

$$e = D(x) \in \mathbb{R}^L$$

Feature Attributions (Saliency Maps)

Slides and Resources

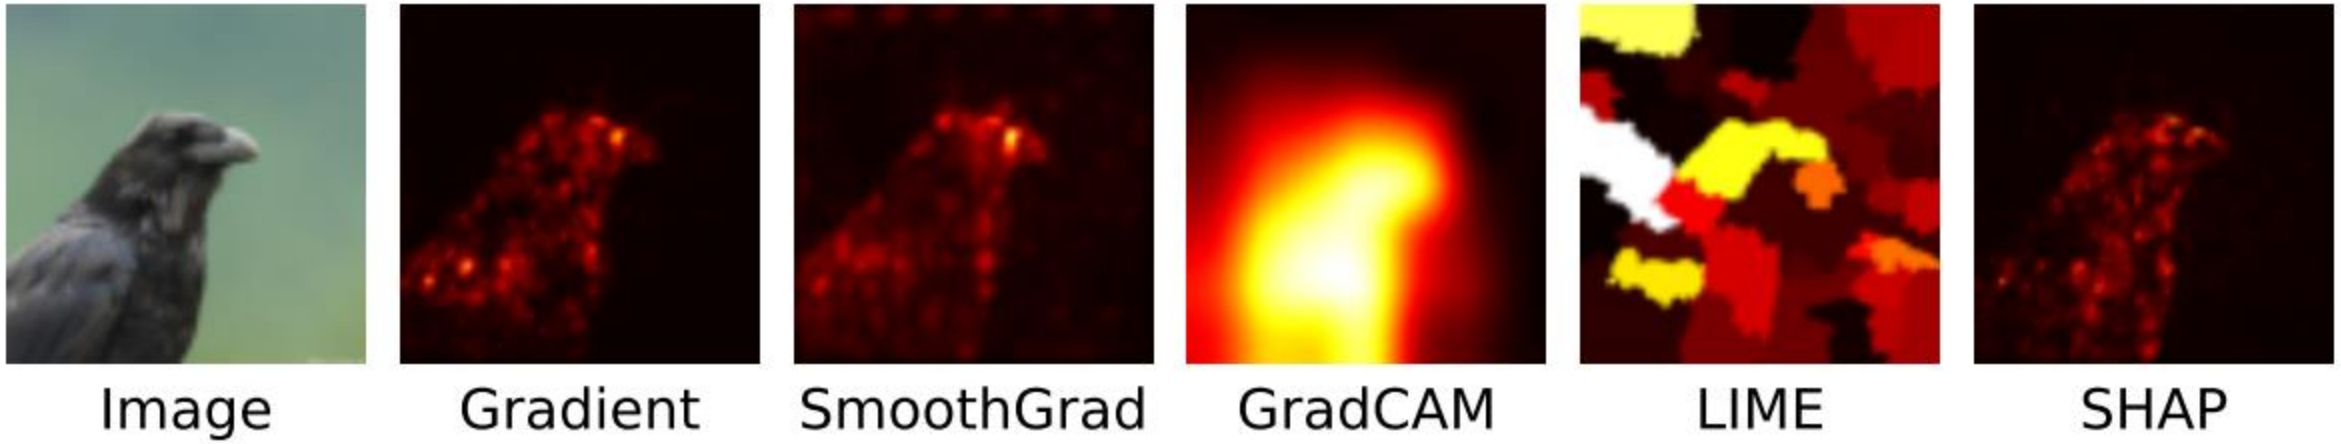

$$e = D(x) \in \mathbb{R}^L$$

Evaluating Explanations

Slides and Resources

Problem: don't know how models work

Evaluating Explanations

Slides and Resources

Evaluating Explanations

Slides and Resources

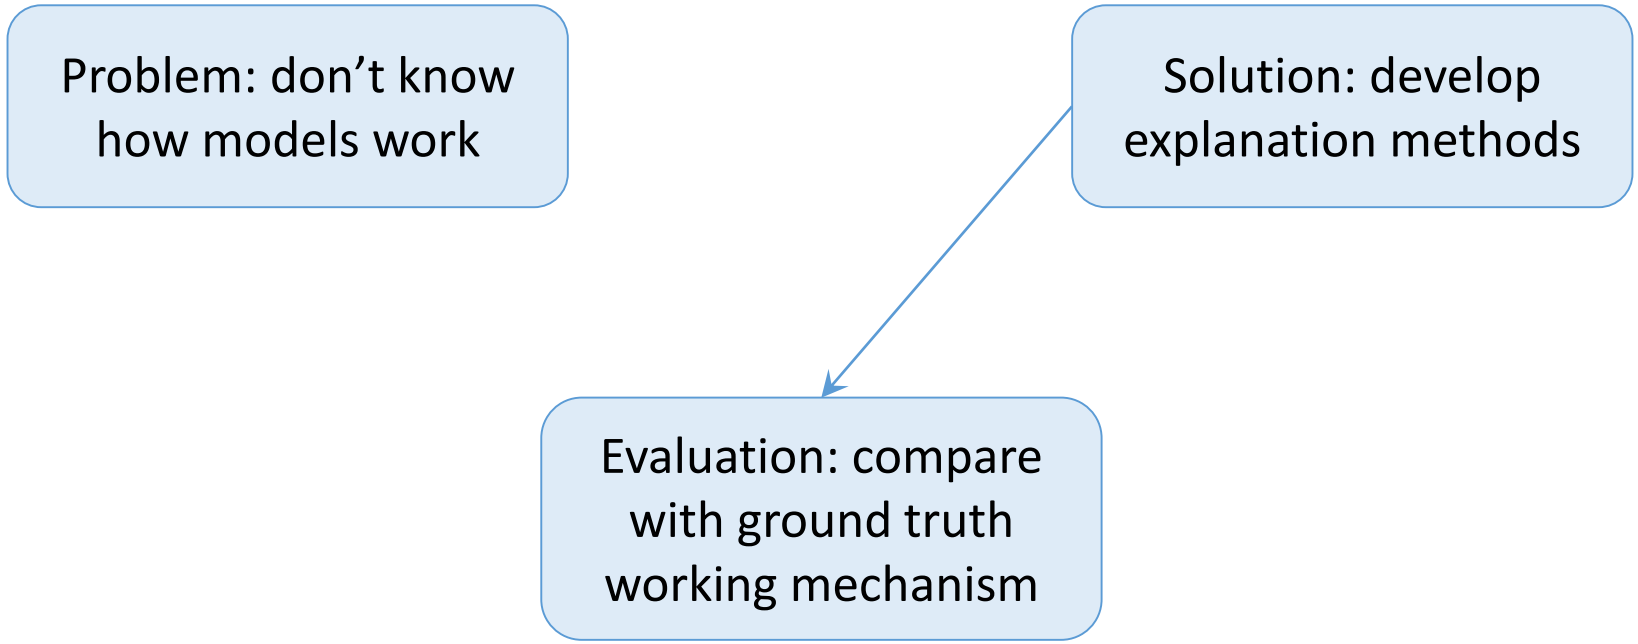

Evaluating Explanations

Slides and Resources

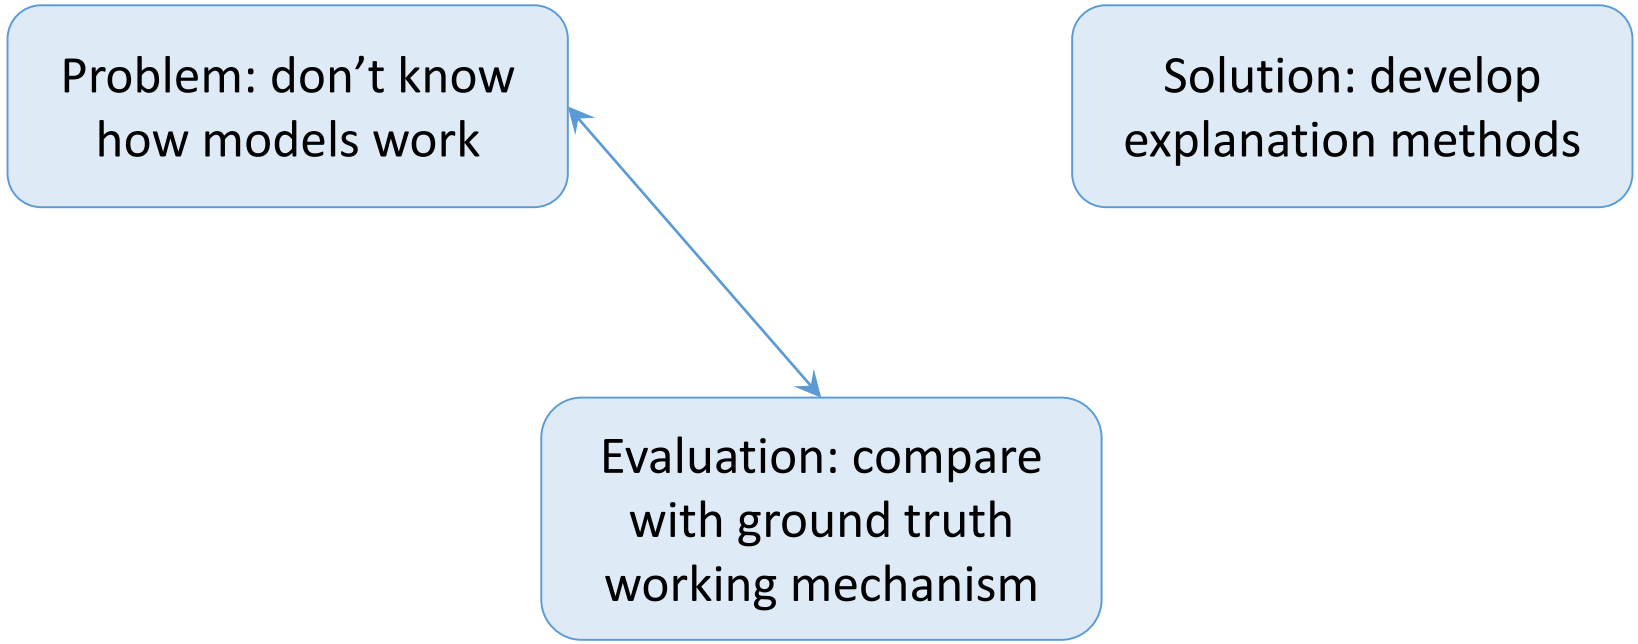

Evaluating Explanations

Slides and Resources

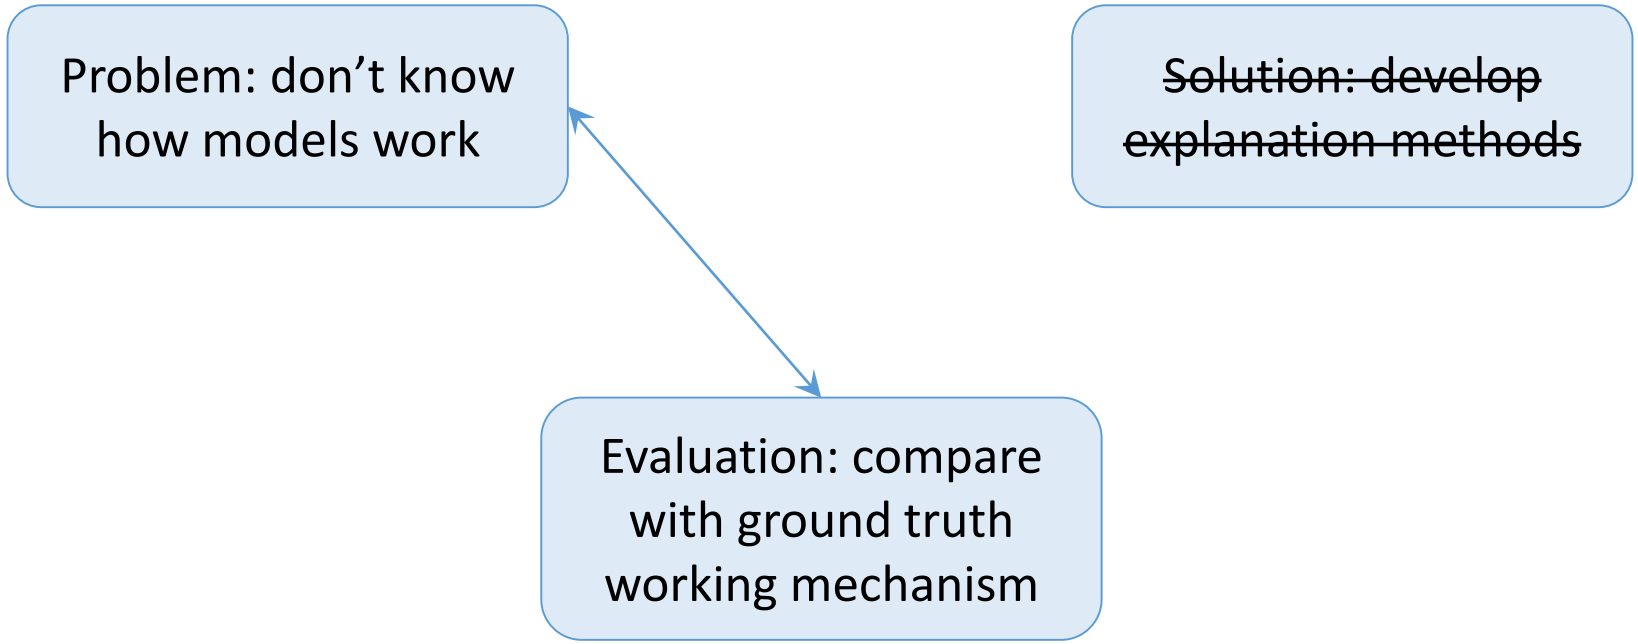

Proxy Metrics for Explanation Quality

Slides and Resources

- Feature importance ⟺ model prediction change with feature removal

Proxy Metrics for Explanation Quality

Slides and Resources

- Feature importance ⟺ model prediction change with feature removal

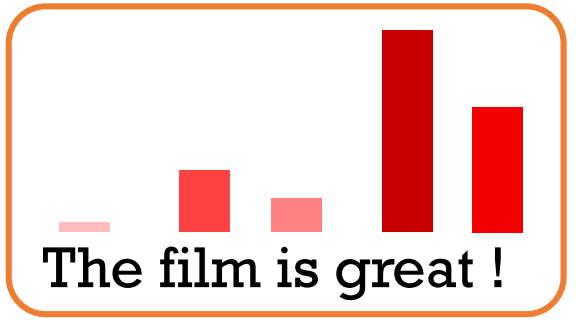

Gradient

Proxy Metrics for Explanation Quality

Slides and Resources

- Feature importance ⟺ model prediction change with feature removal

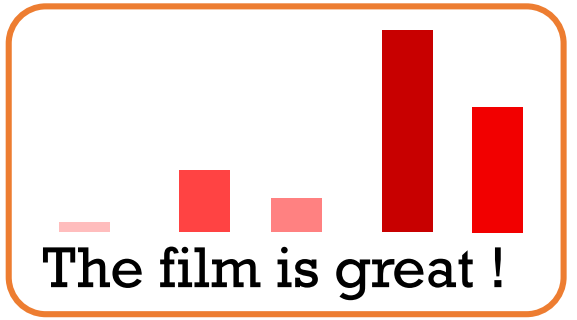

Gradient

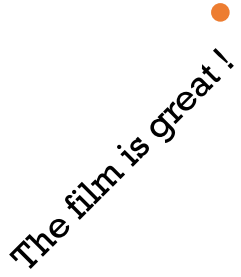

Proxy Metrics for Explanation Quality

Slides and Resources

- Feature importance ⟺ model prediction change with feature removal

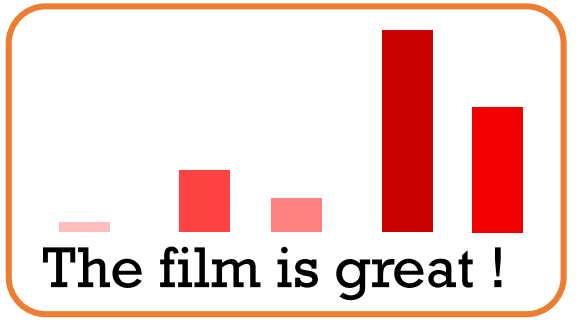

Gradient

Proxy Metrics for Explanation Quality

Slides and Resources

- Feature importance ⟺ model prediction change with feature removal

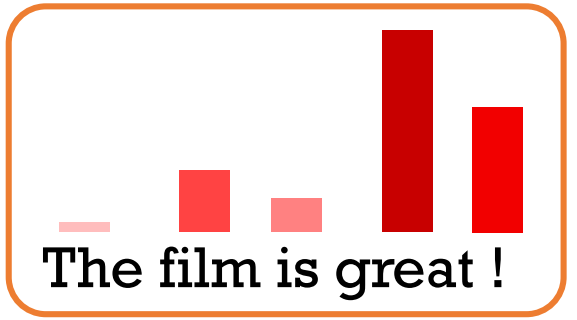

Gradient

Proxy Metrics for Explanation Quality

Slides and Resources

- Feature importance ⟺ model prediction change with feature removal

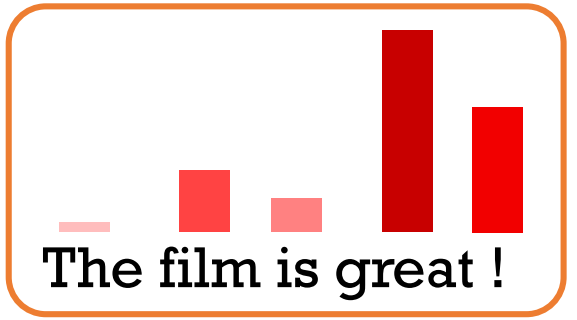

Gradient

Proxy Metrics for Explanation Quality

Slides and Resources

- Feature importance ⟺ model prediction change with feature removal

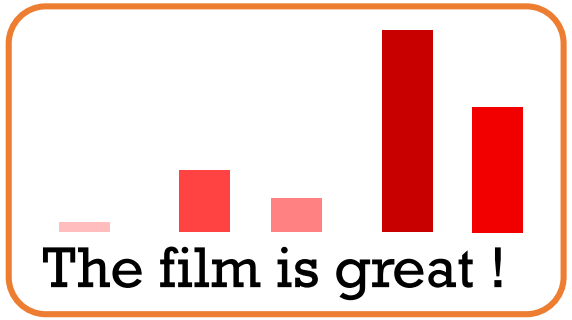

Gradient

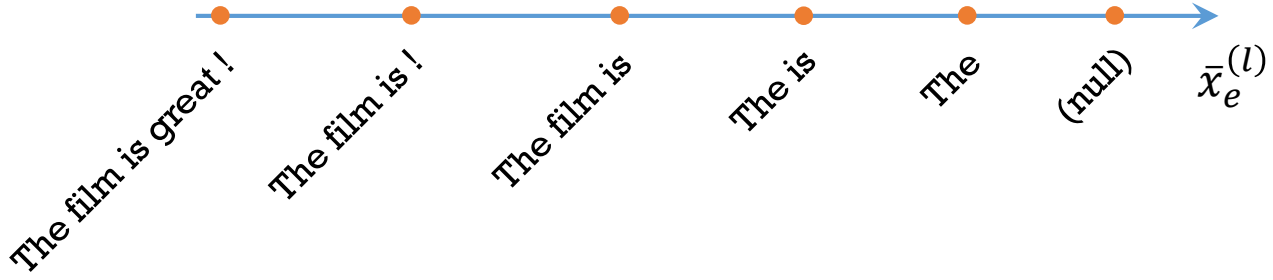

Proxy Metrics for Explanation Quality

Slides and Resources

- Feature importance ⟺ model prediction change with feature removal

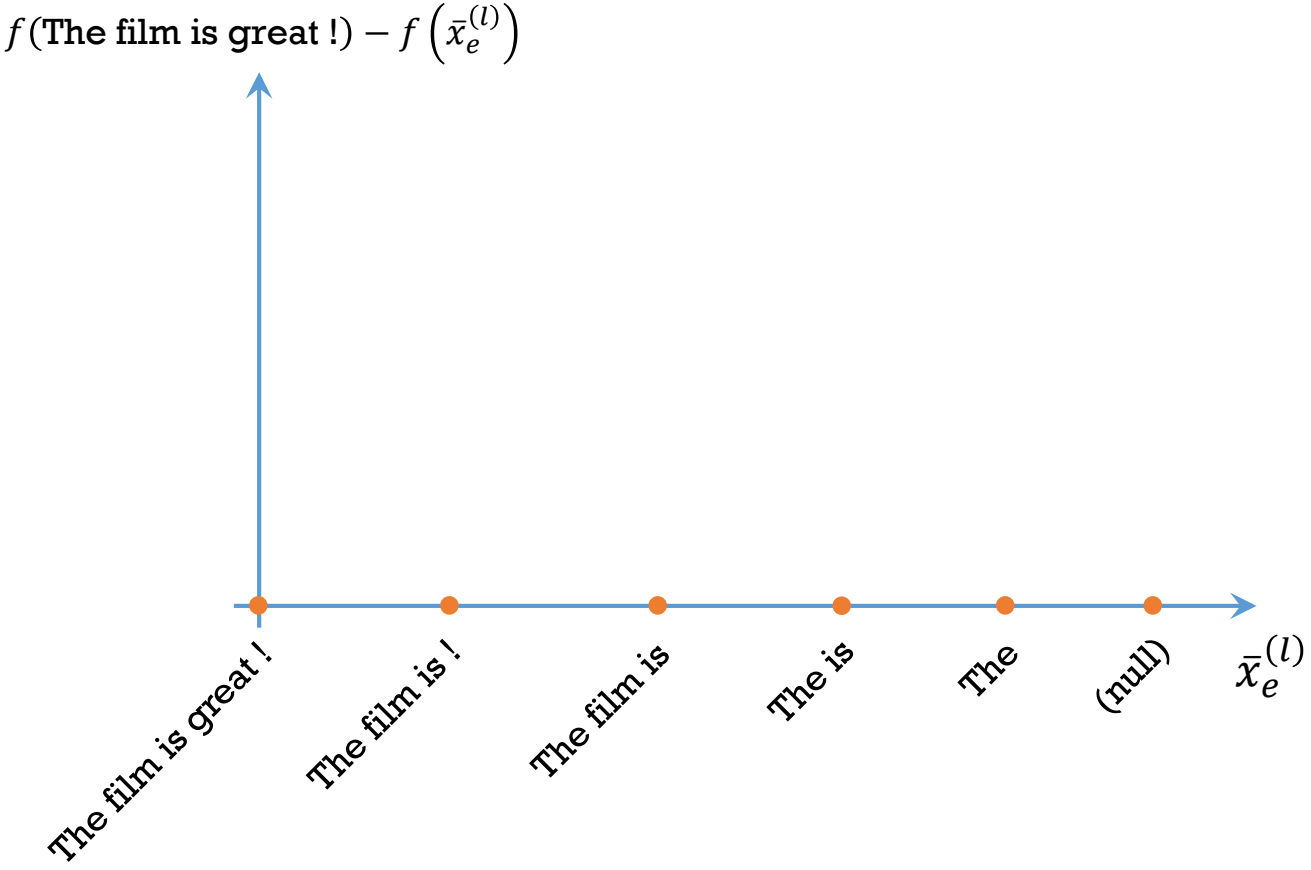

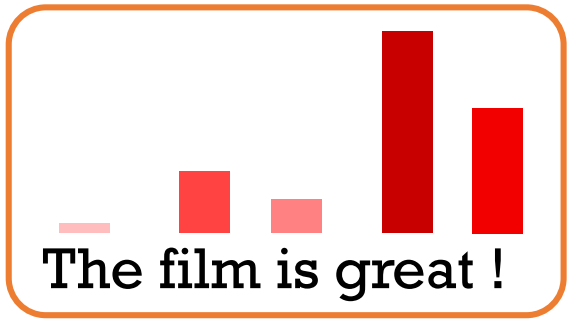

Gradient

Proxy Metrics for Explanation Quality

Slides and Resources

- Feature importance ⟺ model prediction change with feature removal

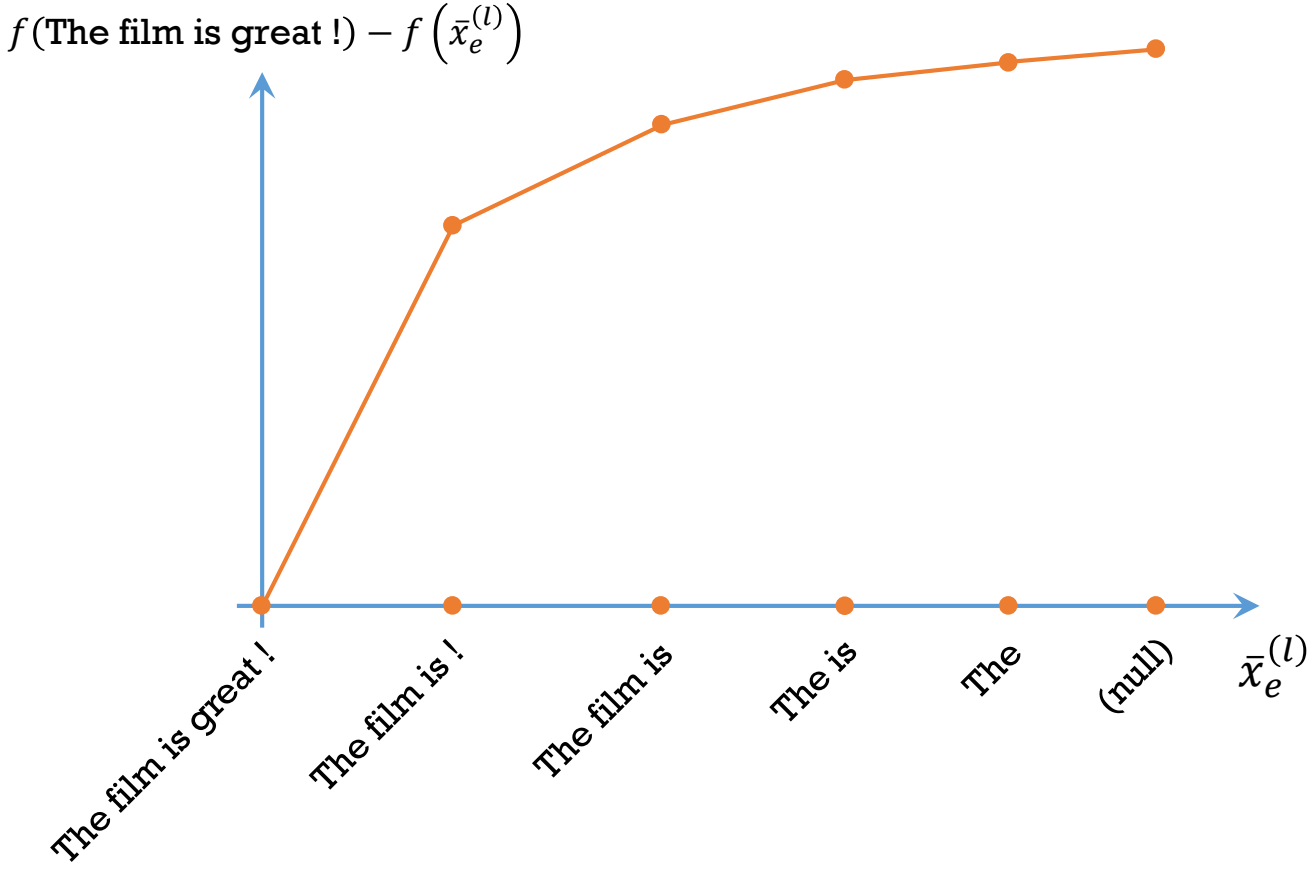

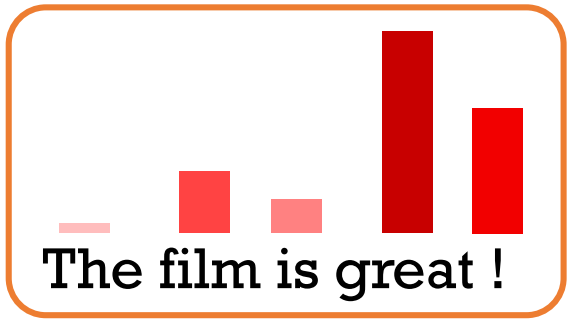

Gradient

Proxy Metrics for Explanation Quality

Slides and Resources

- Feature importance ⟺ model prediction change with feature removal

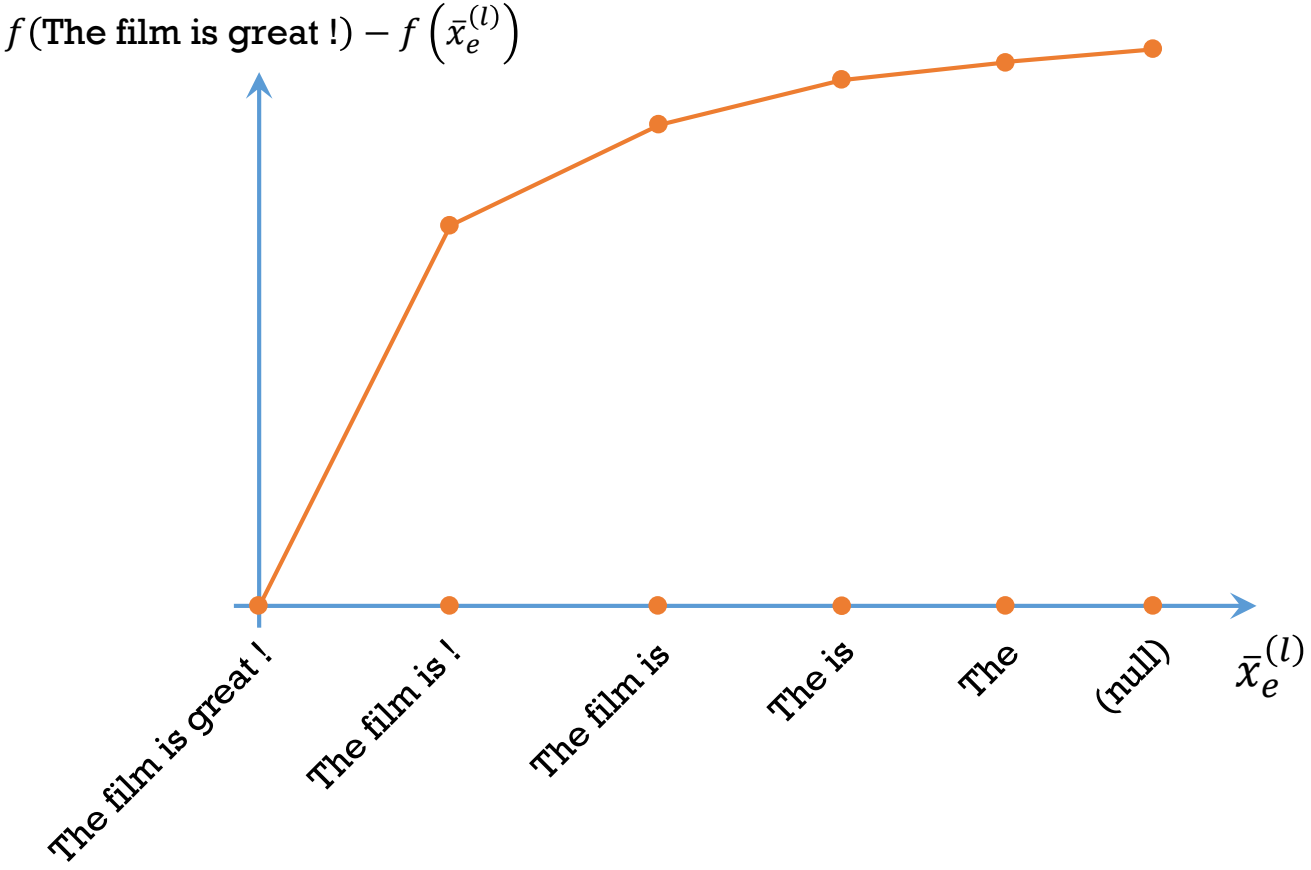

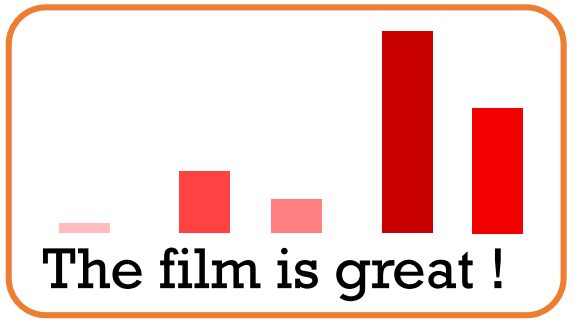

Gradient

Comprehensiveness

$$\kappa(x, e) = \frac{1}{L+1} \sum_{l=0}^{L} f(x) - f\left(\bar{x}_e^{(l)}\right)$$

$\bar{x}^{(l)}$: input x with l most important features removed according to e

$e \in \mathbb{R}^D$ for D-dim input

Proxy Metrics for Explanation Quality

Slides and Resources

- Feature importance ⟺ model prediction change with feature removal

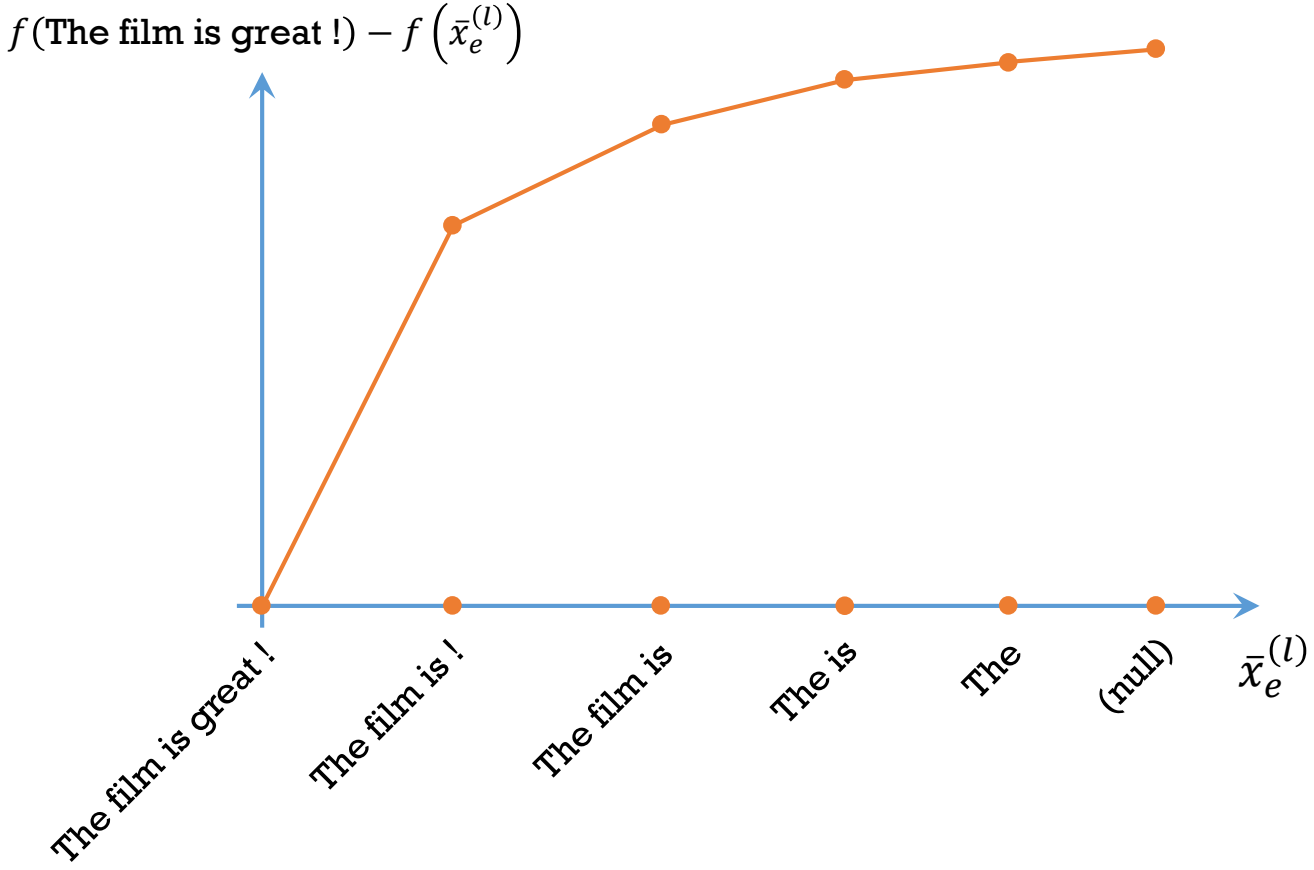

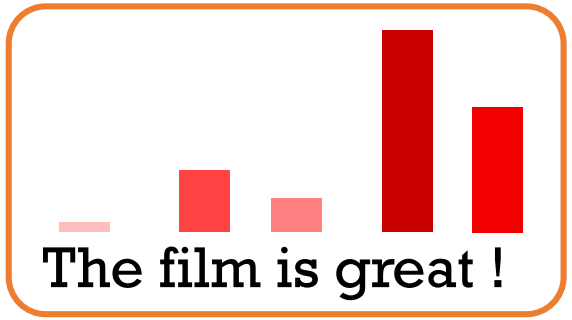

Gradient

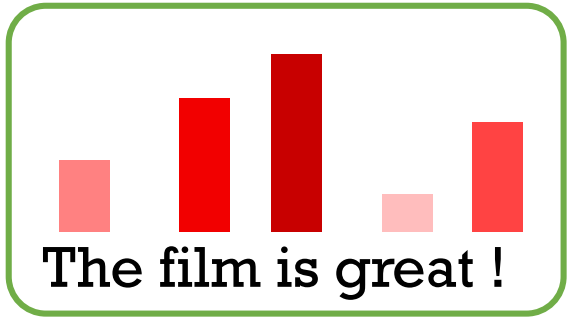

LIME

Proxy Metrics for Explanation Quality

Slides and Resources

- Feature importance ⟺ model prediction change with feature removal

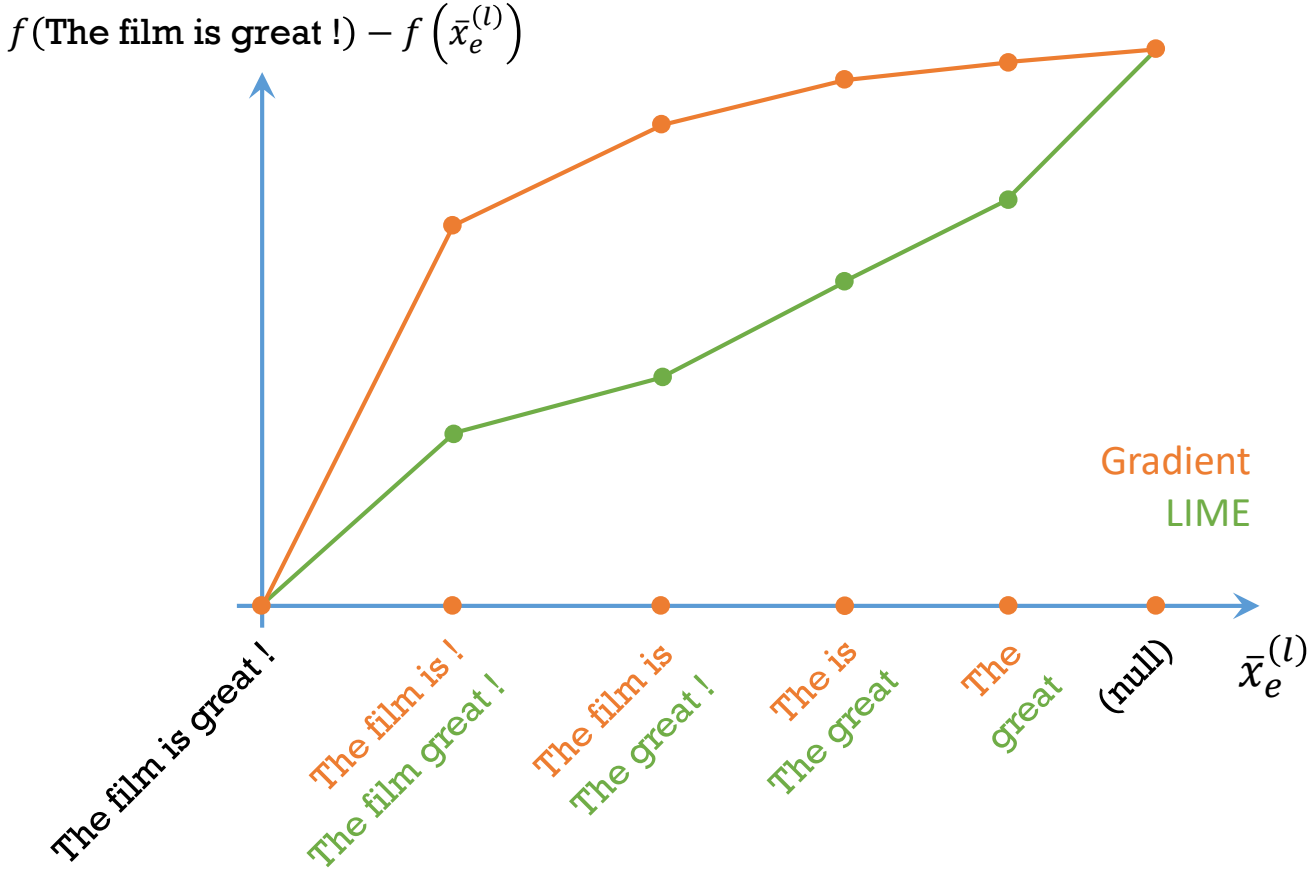

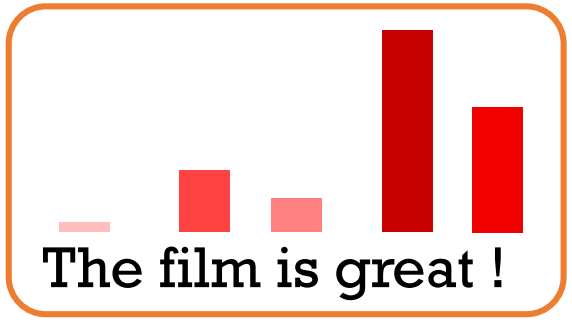

Gradient

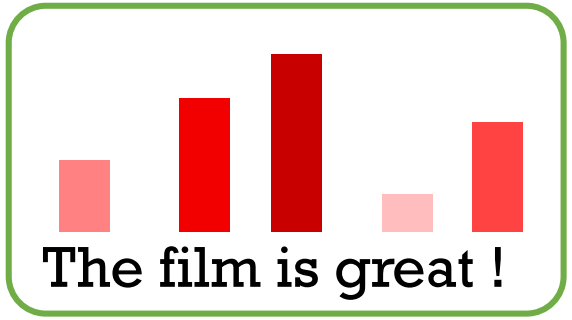

LIME

Common Evaluation Metrics

Slides and Resources

- Comprehensiveness and sufficiency
  - Also known as deletion and insertion metrics
  - Comprehensiveness also known as area over perturbation curve (AoPC)

Common Evaluation Metrics

Slides and Resources

- Comprehensiveness and sufficiency
  - Also known as deletion and insertion metrics
  - Comprehensiveness also known as area over perturbation curve (AoPC)
- Decision flip rate under most important feature removal
- Number of removals required for decision flip
- Prediction change rank correlation
- Etc.

Common Evaluation Metrics

Slides and Resources

- Comprehensiveness and sufficiency
  - Also known as deletion and insertion metrics
  - Comprehensiveness also known as area over perturbation curve (AoPC)
- Decision flip rate under most important feature removal
- Number of removals required for decision flip
- Prediction change rank correlation
- Etc.

Definition: $e = D(x) \in \mathbb{R}^D$

Evaluation: $q = E(x, e) \in \mathbb{R}$

Evaluating Explanations

Slides and Resources

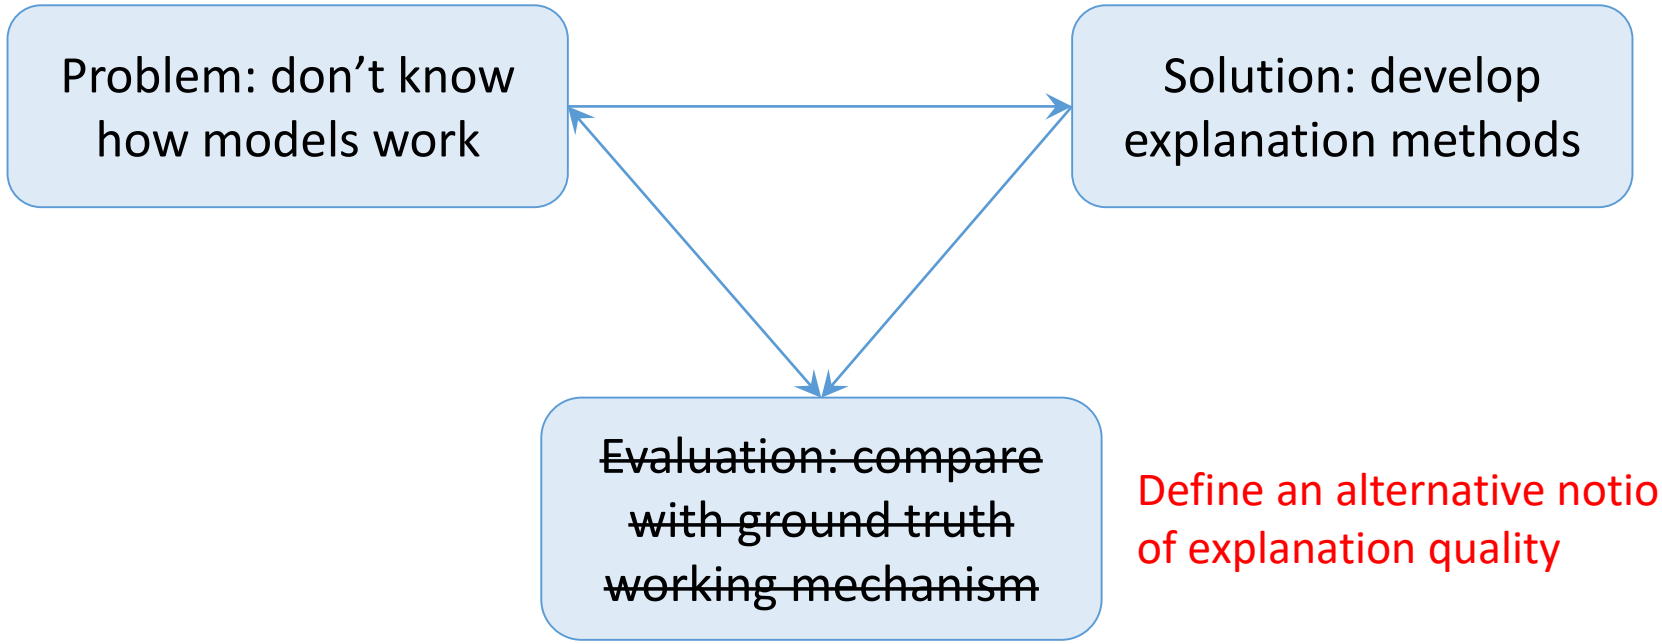

Define an alternative notion of explanation quality

Evaluating Explanations

Slides and Resources

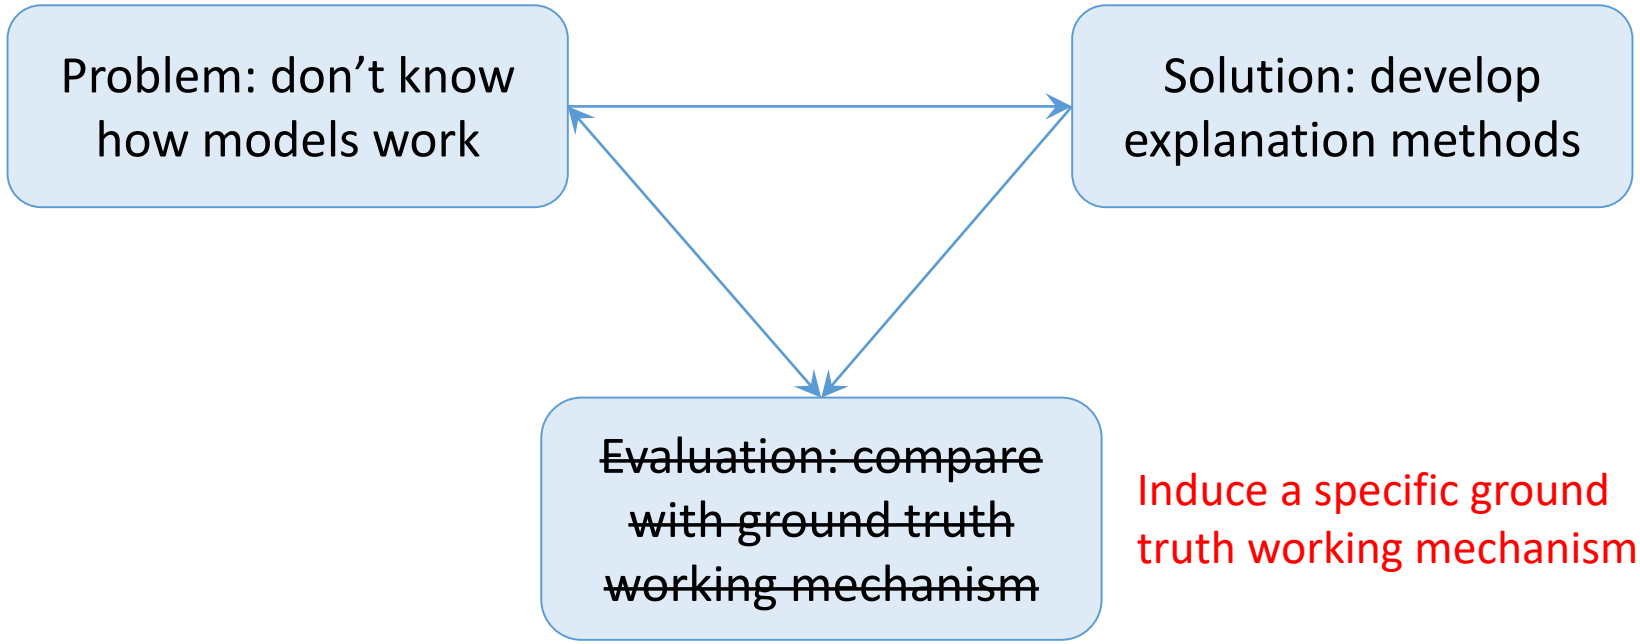

Are Explanations (Necessarily) Correct?

Slides and Resources

If we know that a specific feature is crucial to the model prediction, can feature attribution explanations identify its importance?

Dataset Modification

Slides and Resources

X

Y

→ Crow

→ Crow

→ Sparrow

X: original input image

Y: original output label

Dataset Modification

Slides and Resources

X

Y ✗ $\hat{Y} \sim \text{Ber}(0.5)$


Crow


Crow

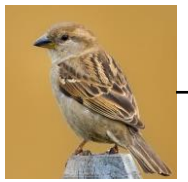
Sparrow


Crow


Sparrow


Crow

X: original input image

Y: original output label

$\hat{Y}$: modified output label

Dataset Modification

Slides and Resources

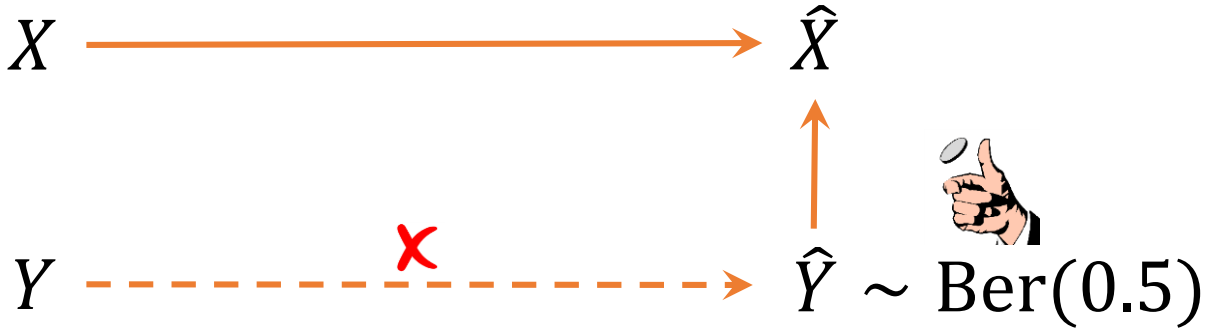


Crow


Crow


Sparrow

Crow

Sparrow

Crow

X: original input image

Y: original output label

$\hat{Y}$: modified output label

$\hat{X}$: modified input image

Dataset Modification

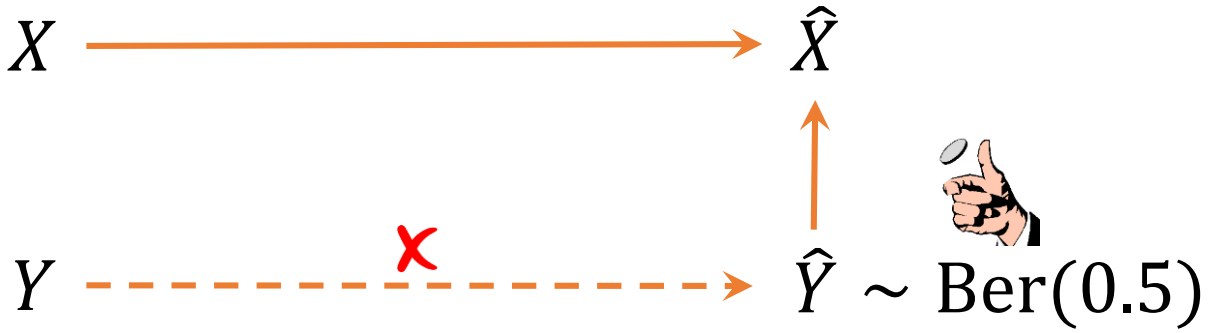

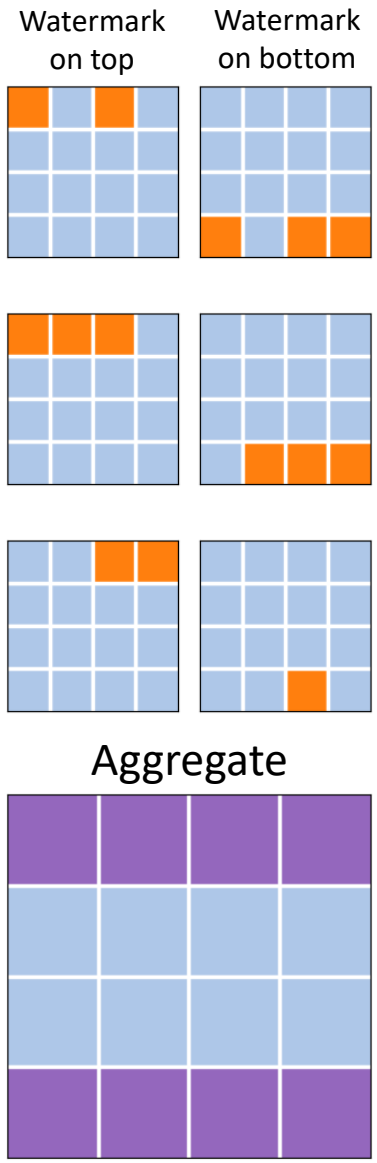

Slides and Resources

Dataset Modification

Slides and Resources

Dataset Modification

Slides and Resources

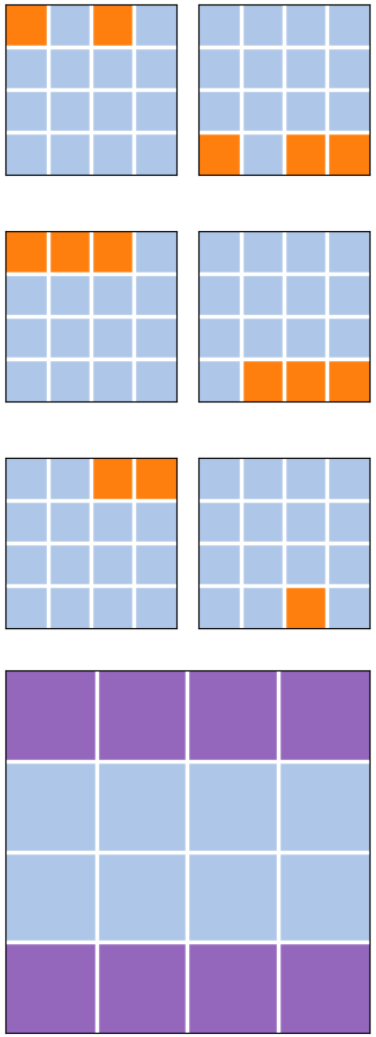

$$\% \text{ Attribution} = \frac{\sum_{\blacksquare} A_{\blacksquare}}{\sum_{\blacksquare} A_{\blacksquare} + \sum_{\blacksquare} A_{\blacksquare}}$$

Evaluating Image Saliency Maps

Slides and Resources

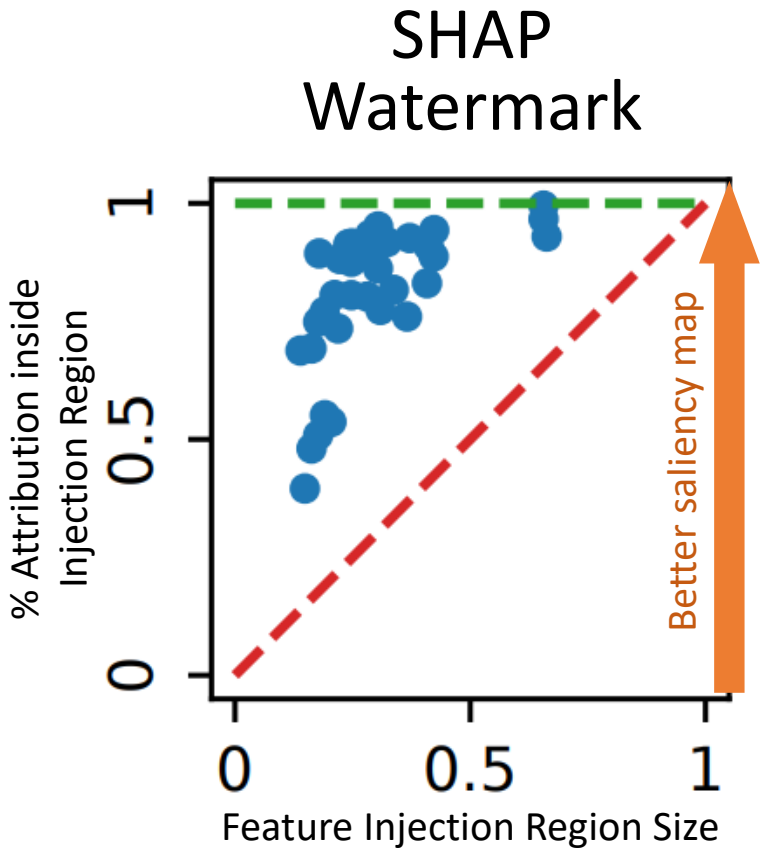

Random saliency map

Optimal saliency map

Evaluating Image Saliency Maps

Slides and Resources

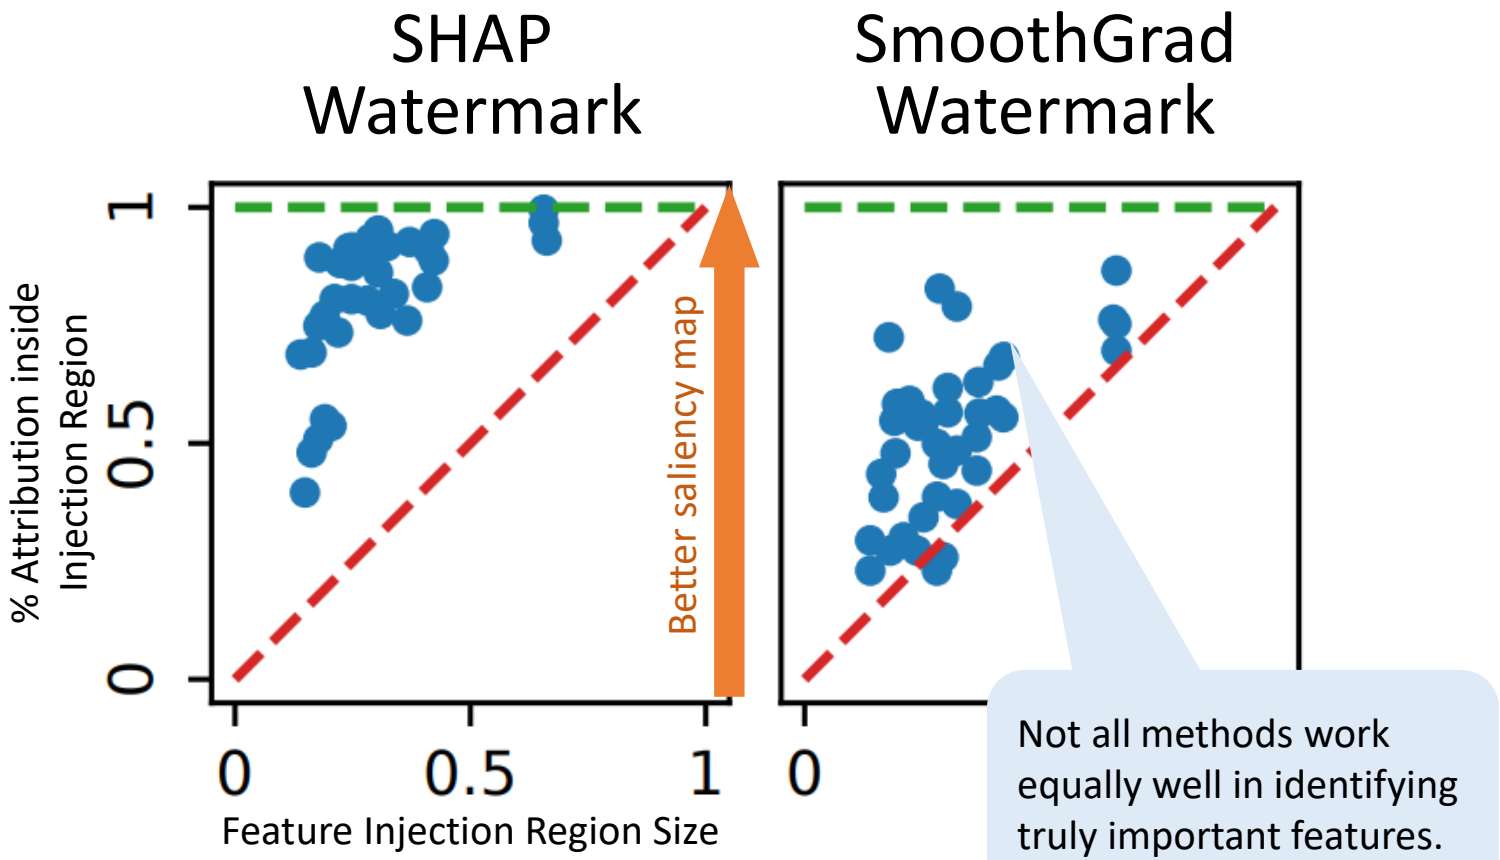

Evaluating Image Saliency Maps

Slides and Resources

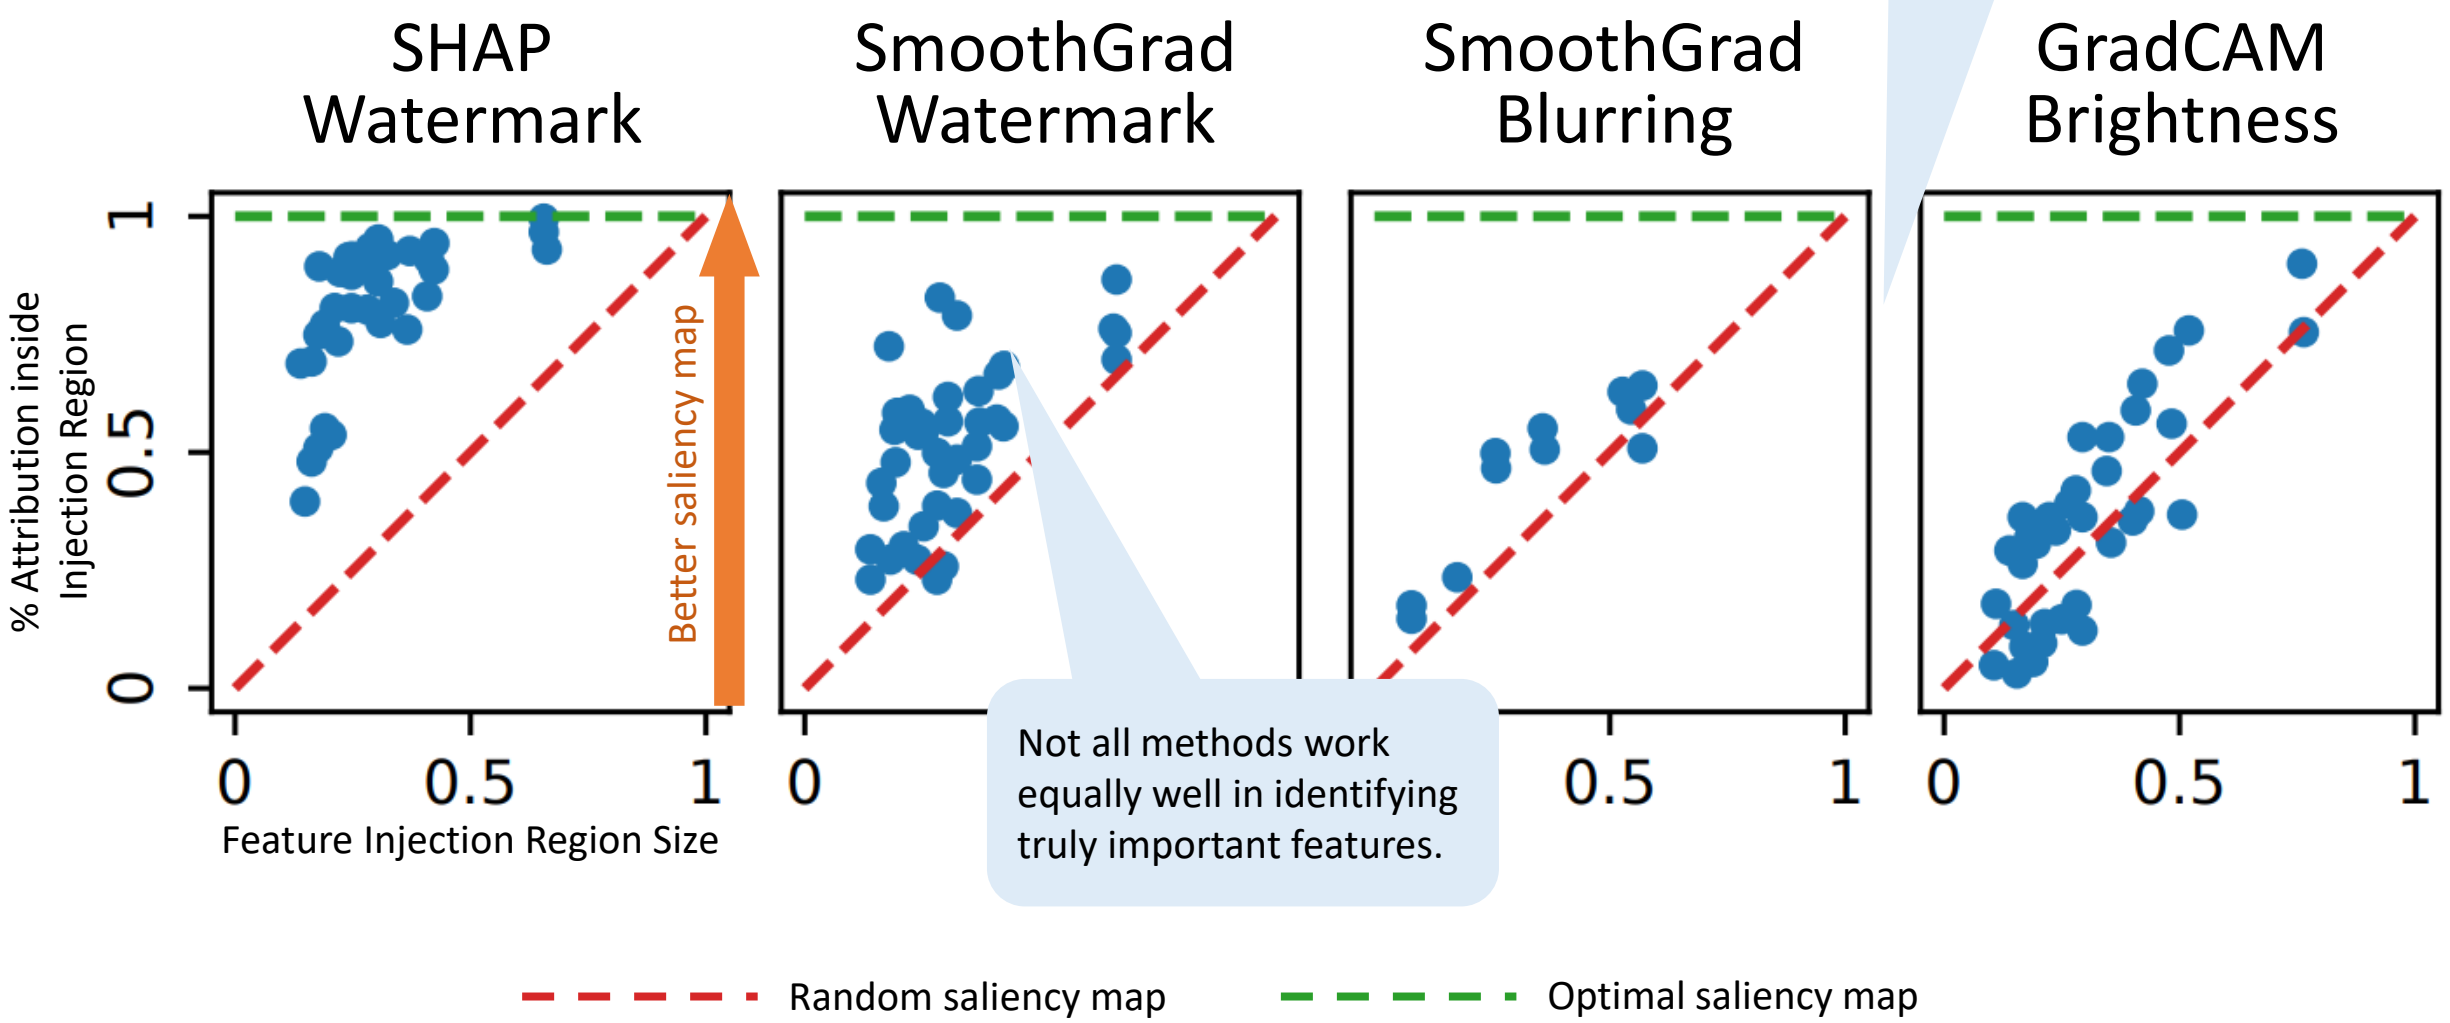

Practical Implications

Slides and Resources

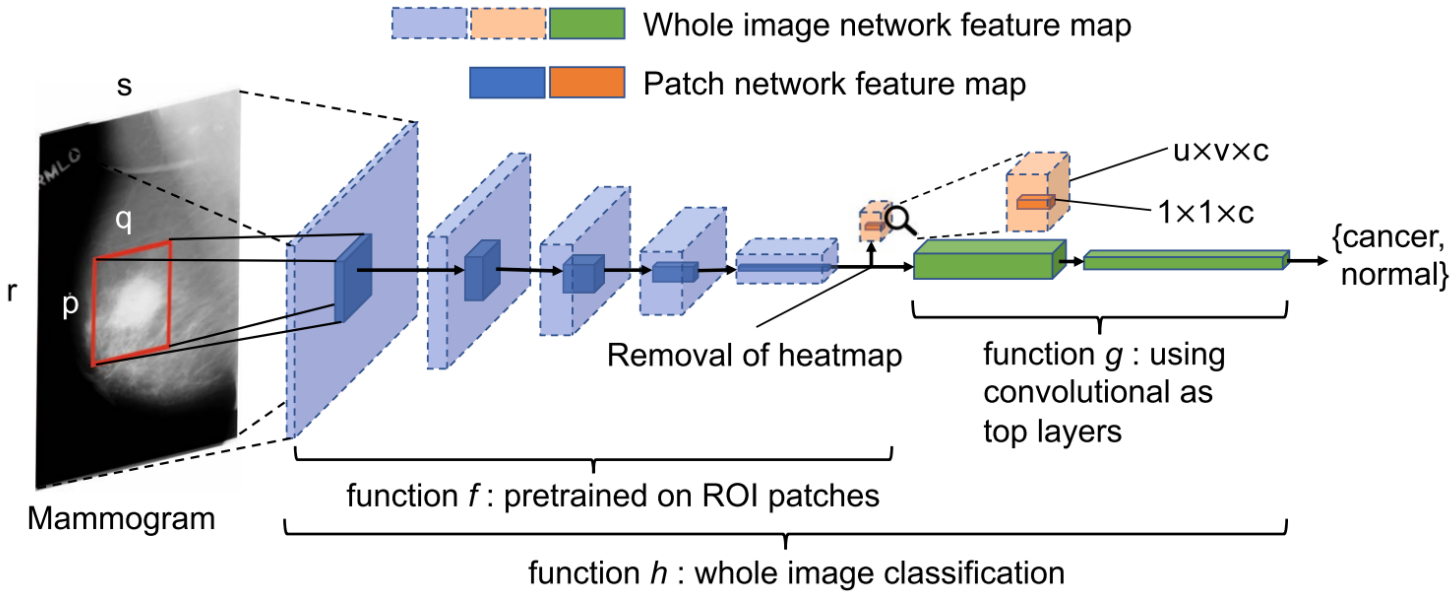

Practical Implications

Slides and Resources

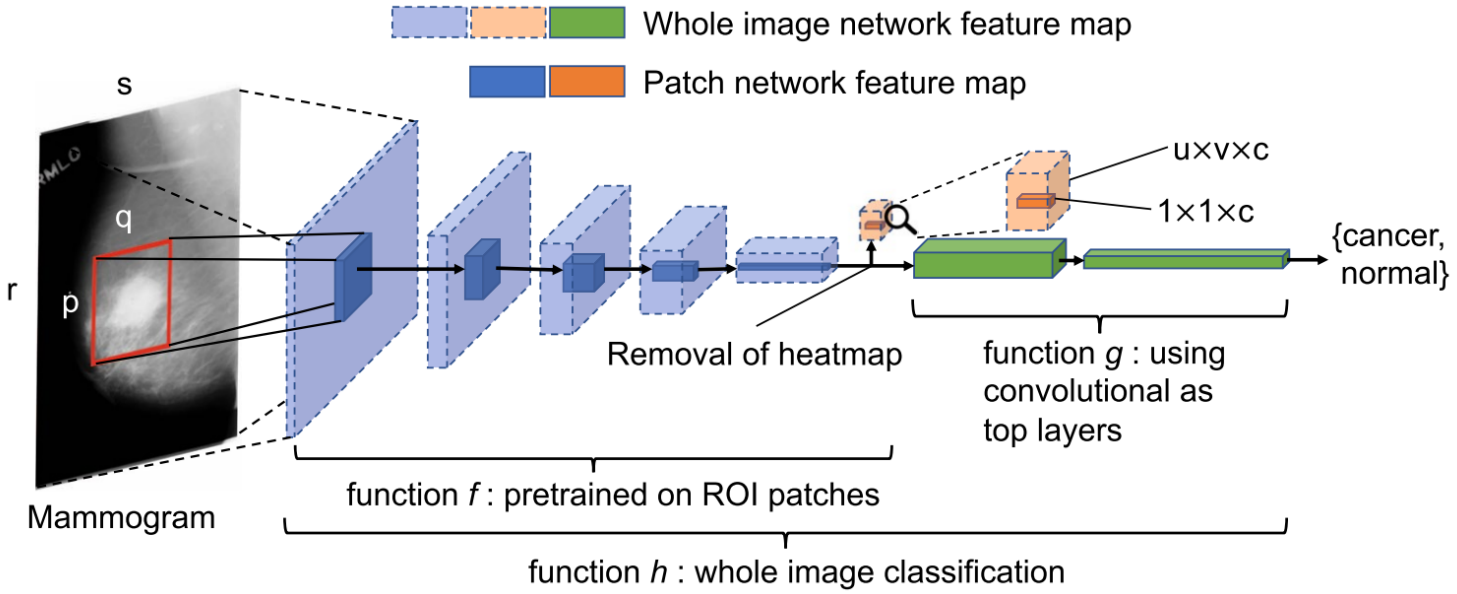

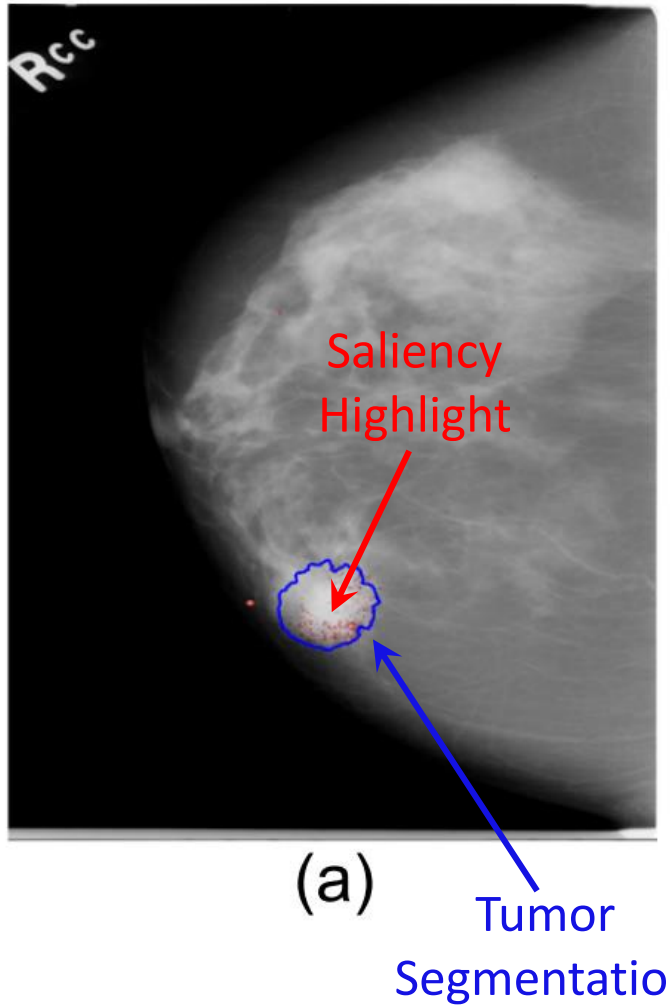

"A saliency map illustrates which area of the input image is considered to be responsible for the cancer […] Figure 4a […] shows that the image classifier was able to correctly locate the cancerous region on which its decision was based."

(Shen et al., *Scientific Reports*, 2019)

Practical Implications

Slides and Resources

Known spurious correlation

Some explainers don't work

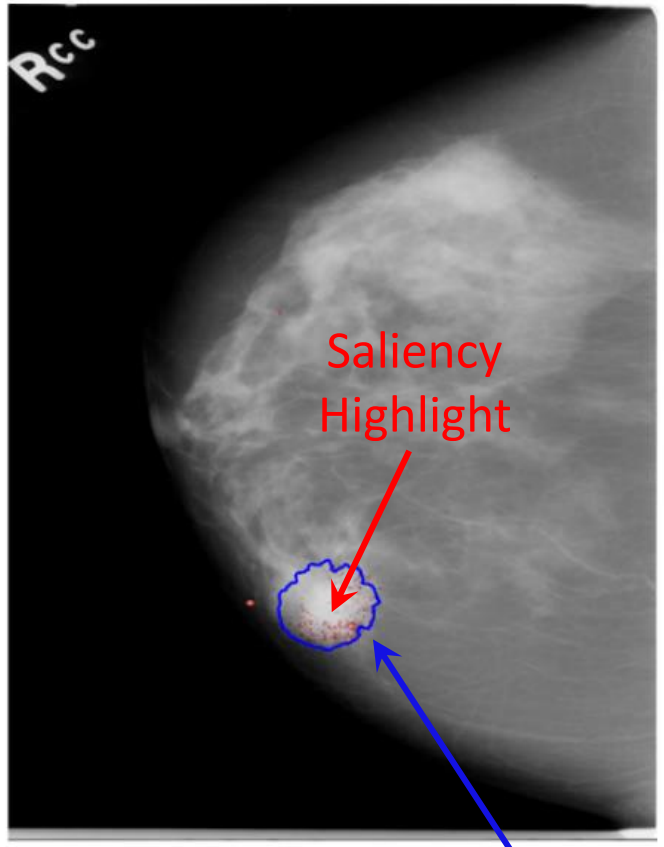

"A saliency [...] rea of the input image is considered to be responsible for the cancer [...] Figure 4a [...] shows that the image classifier was able to correctly locate the cancerous region on which its decision was based."

(Shen et al., *Scientific Reports*, 2019)

Practical Implications

Slides and Resources

Known spurious correlation

Some explainers don't work

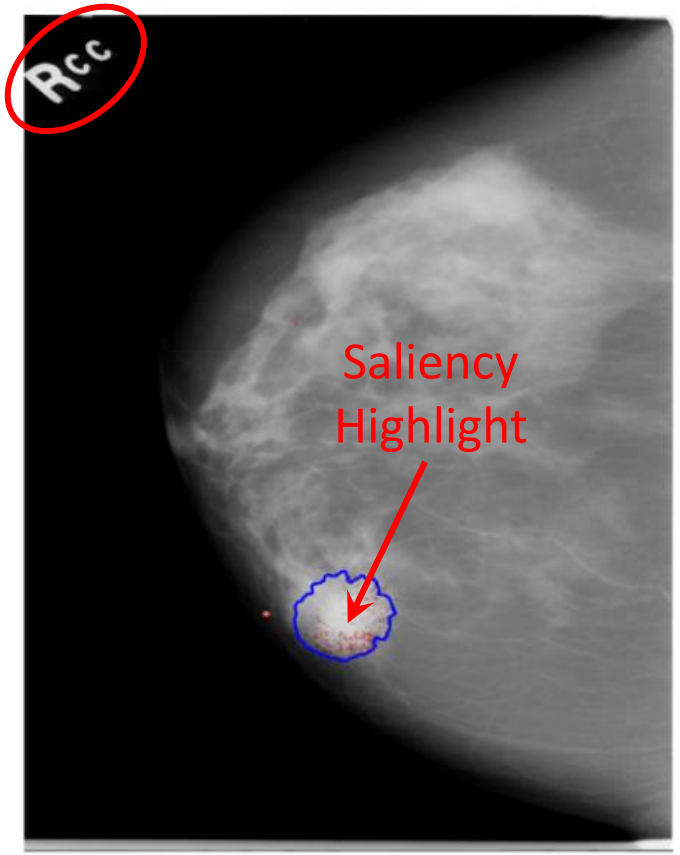

Unknown spurious correlation

Can we trust these explainers?

"A saliency [...] rea of the input image is considered to be responsible for the cancer [...] Figure 4a [...] shows that the image classifier was able to correctly locate the cancerous region on which its decision was based."

(Shen et al., *Scientific Reports*, 2019)

Explanation Understandability


Slides and Resources

Do the explanations correctly explain the model prediction logic?

Explanation Understandability

Slides and Resources

Do the explanations correctly explain the model prediction logic?

Do people correctly understand the model prediction logic from the explanations?

Correct but Not Understandable Explanations

Slides and Resources

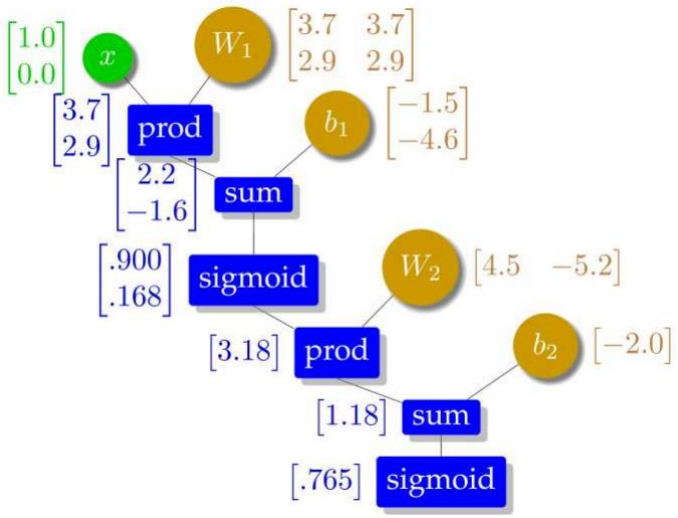

Computation trace: a fully correct explanation for the prediction

Correct but Not Understandable Explanations

Slides and Resources

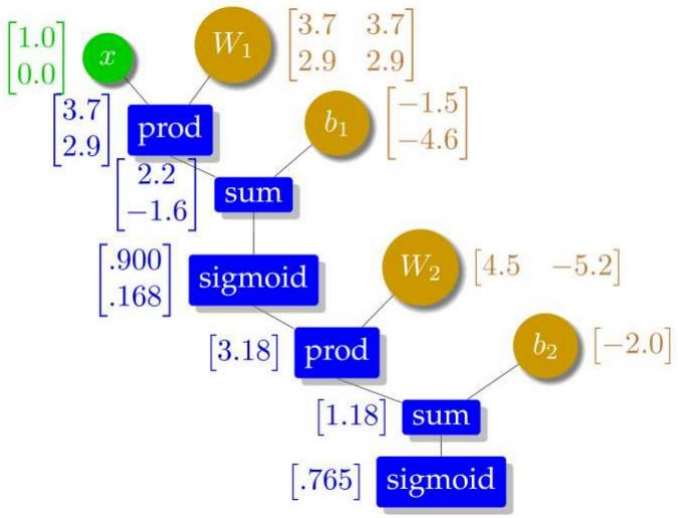

Computation trace: a fully correct but totally not understandable explanation for the prediction

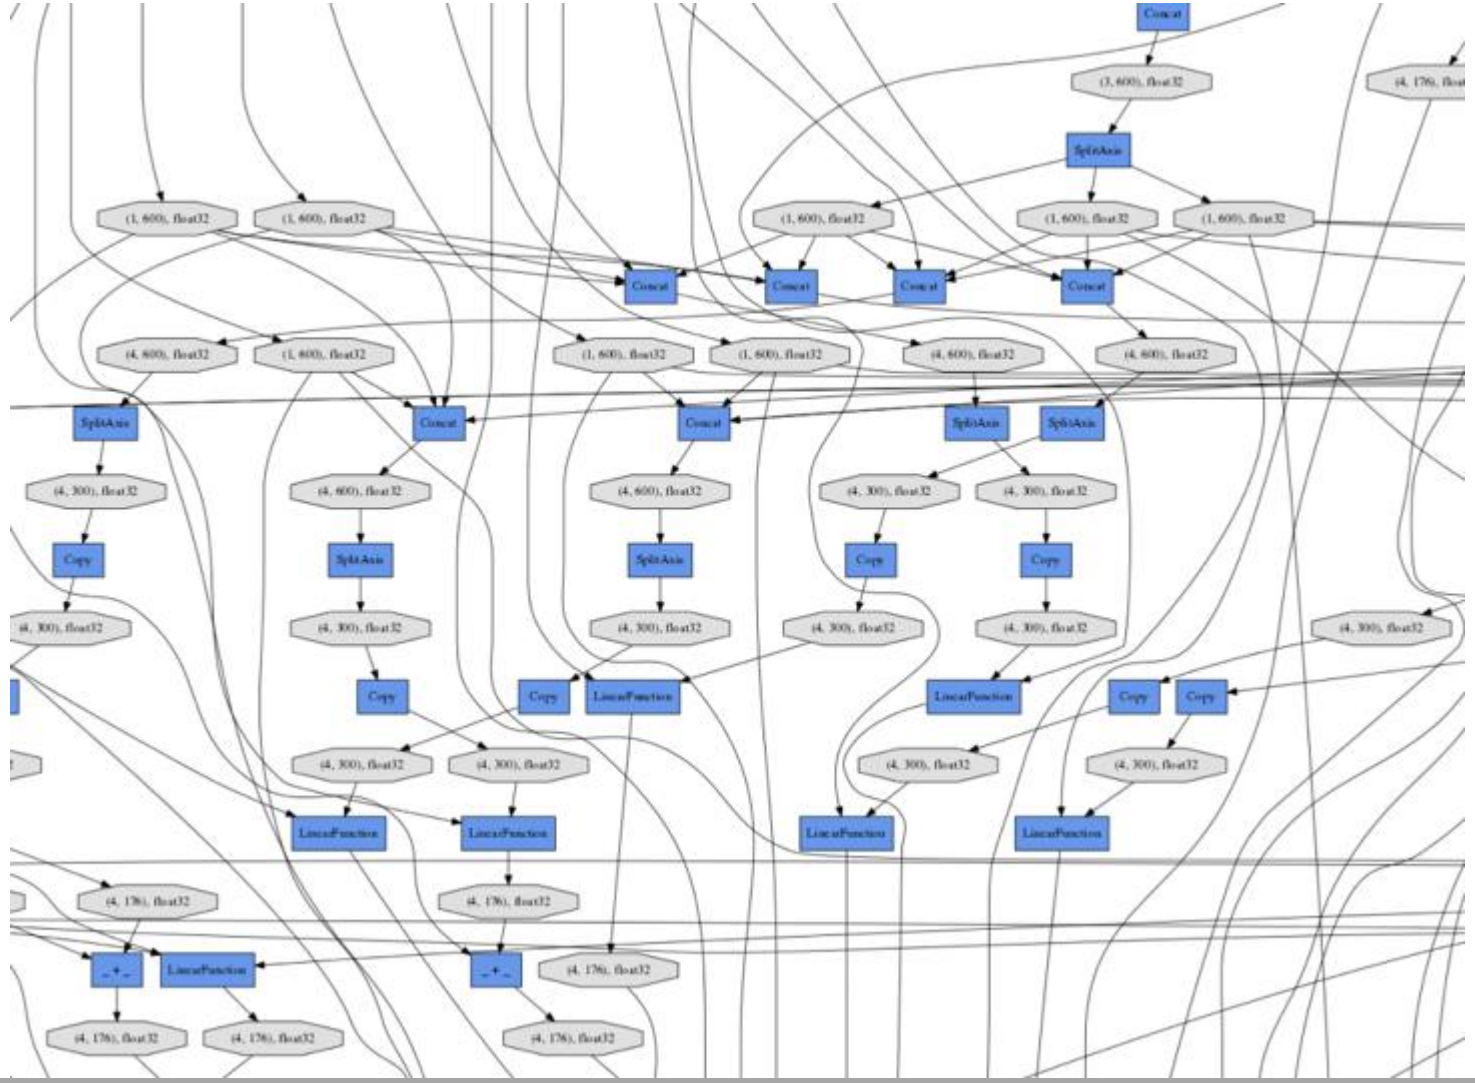

Understanding a Sentiment Classifier

Slides and Resources

Data Instance

Input: As shaky as the plot is , Kaufman 's script is still memorable for some great one-liners .
Label: Positive
Model Prediction: Positive

Local Explanation

SHAP score

memorable:	0.48
great:	0.37
for:	-0.02
one-liners:	-0.14
shaky:	-0.39

Understanding a Sentiment Classifier

Slides and Resources

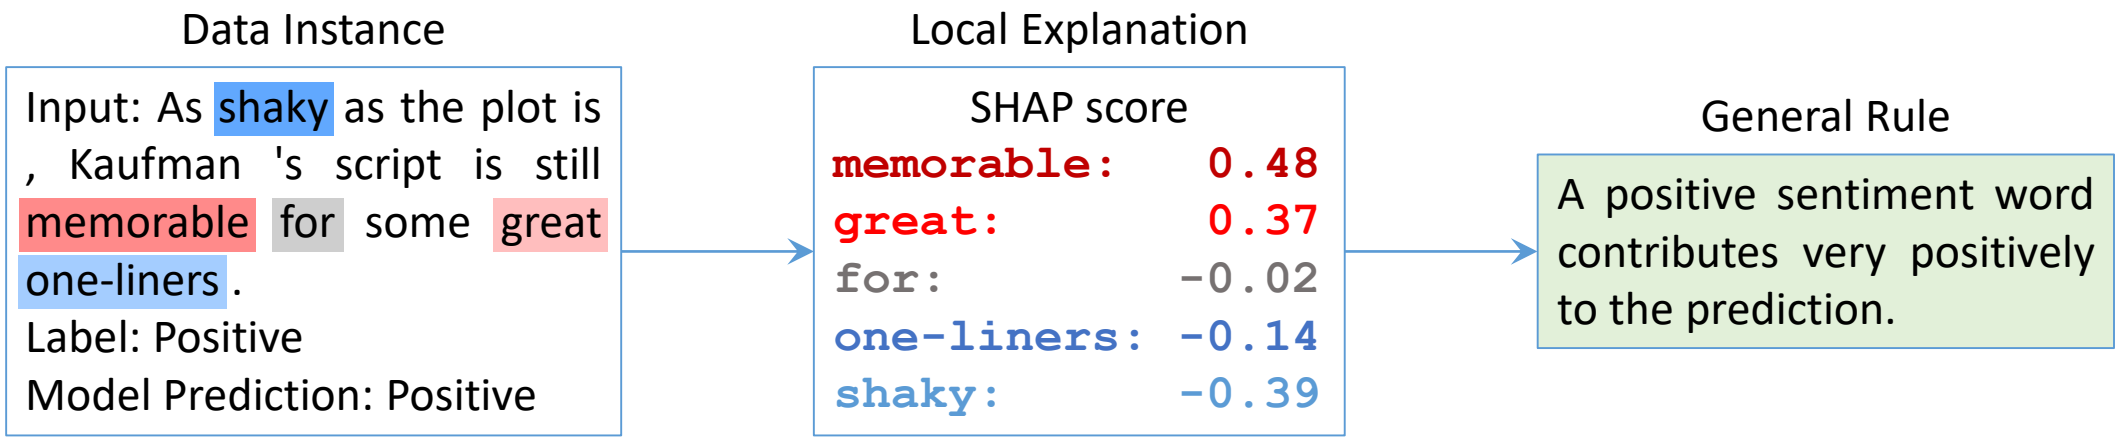

Understanding a Sentiment Classifier


Slides and Resources

As shaky as the plot is, Kaufman's script is still **memorable** for some **great** one-liners.

```
memorable:  0.48
great:      0.37
```

Label = Positive
Prediction = Positive

A positive sentiment word contributes very positively to the prediction

Understanding a Sentiment Classifier


Slides and Resources

As shaky as the plot is, Kaufman's script is still **memorable** for some **great** one-liners.

```
memorable:  0.48
great:      0.37
```

Label = Positive
Prediction = Positive

Just because it really happened to you, honey, does n't mean that it's **attractive** to anyone else.

```
attractive:  0.00
```

Label = Negative
Prediction = Negative

A positive sentiment word contributes very positively to the prediction

Supporting instance

Opposing instance

Understanding a Sentiment Classifier

Slides and Resources

As shaky as the plot is, Kaufman's script is still **memorable** for some **great** one-liners.

```
memorable:  0.48
great:      0.37
```

Label = Positive
Prediction = Positive

Just because it really happened to you, honey, does n't mean that it's **attractive** to anyone else.

```
attractive:  0.00
```

Label = Negative
Prediction = Negative

A positive sentiment word contributes very positively to the prediction…

unless it is negated

Supporting instance

Opposing instance

Understanding a Sentiment Classifier


Slides and Resources

As shaky as the plot is, Kaufman's script is still **memorable** for some **great** one-liners.

```
memorable:  0.48
great:      0.37
```

Label = Positive
Prediction = Positive

Just because it really happened to you, honey, does n't mean that it's **attractive** to anyone else.

```
attractive:  0.00
```

Label = Negative
Prediction = Negative

A positive sentiment word contributes very positively to the prediction…

unless it is negated

Daring, **mesmerizing** and exceed-ingly hard to forget.

```
Daring:        0.13
mesmerizing:   0.06
```

Label = Positive
Prediction = Positive

Supporting instance

Opposing instance

Understanding a Sentiment Classifier


Slides and Resources

As shaky as the plot is, Kaufman's script is still **memorable** for some **great** one-liners.

```
memorable:  0.48
great:      0.37
```

Label = Positive
Prediction = Positive

Just because it really happened to you, honey, does n't mean that it's **attractive** to anyone else.

```
attractive:  0.00
```

Label = Negative
Prediction = Negative

A positive sentiment word contributes very positively to the prediction…

unless it is negated, or near another positive word

Daring, **mesmerizing** and exceed-ingly hard to forget.

```
Daring:       0.13
mesmerizing:  0.06
```

Label = Positive
Prediction = Positive

Supporting instance

Opposing instance

Understanding a Sentiment Classifier

Slides and Resources

As shaky as the plot is, Kaufman's script is still **memorable** for some **great** one-liners.

```
memorable:  0.48
great:      0.37
```

Label = Positive
Prediction = Positive

Just because it really happened to you, honey, does n't mean that it's **attractive** to anyone else.

```
attractive:  0.00
```

Label = Negative
Prediction = Negative

A positive sentiment word contributes very positively to the prediction…

unless it is negated, or near another positive word

Daring, **mesmerizing** and exceed-ingly hard to forget.

```
Daring:       0.13
mesmerizing:  0.06
```

Label = Positive
Prediction = Positive

Ranks among Willams' **best** screen work.

```
best:  0.05
```

Label = Positive
Prediction = Positive

Supporting instance

Opposing instance

Understanding a Sentiment Classifier

Slides and Resources

As shaky as the plot is, Kaufman's script is still **memorable** for some **great** one-liners.

```
memorable:  0.48
great:      0.37
```

Label = Positive
Prediction = Positive

Just because it really happened to you, honey, does n't mean that it's **attractive** to anyone else.

```
attractive:  0.00
```

Label = Negative
Prediction = Negative

A positive sentiment word sometimes contributes very positively to the prediction…

unless it is negated, or near another positive word

Daring, **mesmerizing** and exceed-ingly hard to forget.

```
Daring:        0.13
mesmerizing:   0.06
```

Label = Positive
Prediction = Positive

Ranks among Willams' **best** screen work.

```
best:  0.05
```

Label = Positive
Prediction = Positive

Supporting instance

Opposing instance

Understanding a Sentiment Classifier

Slides and Resources

As shaky as the plot is, Kaufman's script is still **memorable** for some **great** one-liners.

```
memorable:  0.48
great:      0.37
```

Label = Positive
Prediction = Positive

Just because it really happened to you, honey, does n't mean that it's **attractive** to anyone else.

```
attractive:  0.00
```

Label = Negative
Prediction = Negative

A positive sentiment word sometimes contributes very positively to the prediction…

unless it is negated, or near another positive word

Daring, **mesmerizing** and exceedingly hard to forget.

```
Daring:       0.13
mesmerizing:  0.06
```

Label = Positive
Prediction = Positive

Ranks among Willams' **best** screen work.

```
best:  0.05
```

Label = Positive
Prediction = Positive

(many more counter-examples not shown)

Supporting instance

Opposing instance

Explanation Summary (ExSum)

Slides and Resources

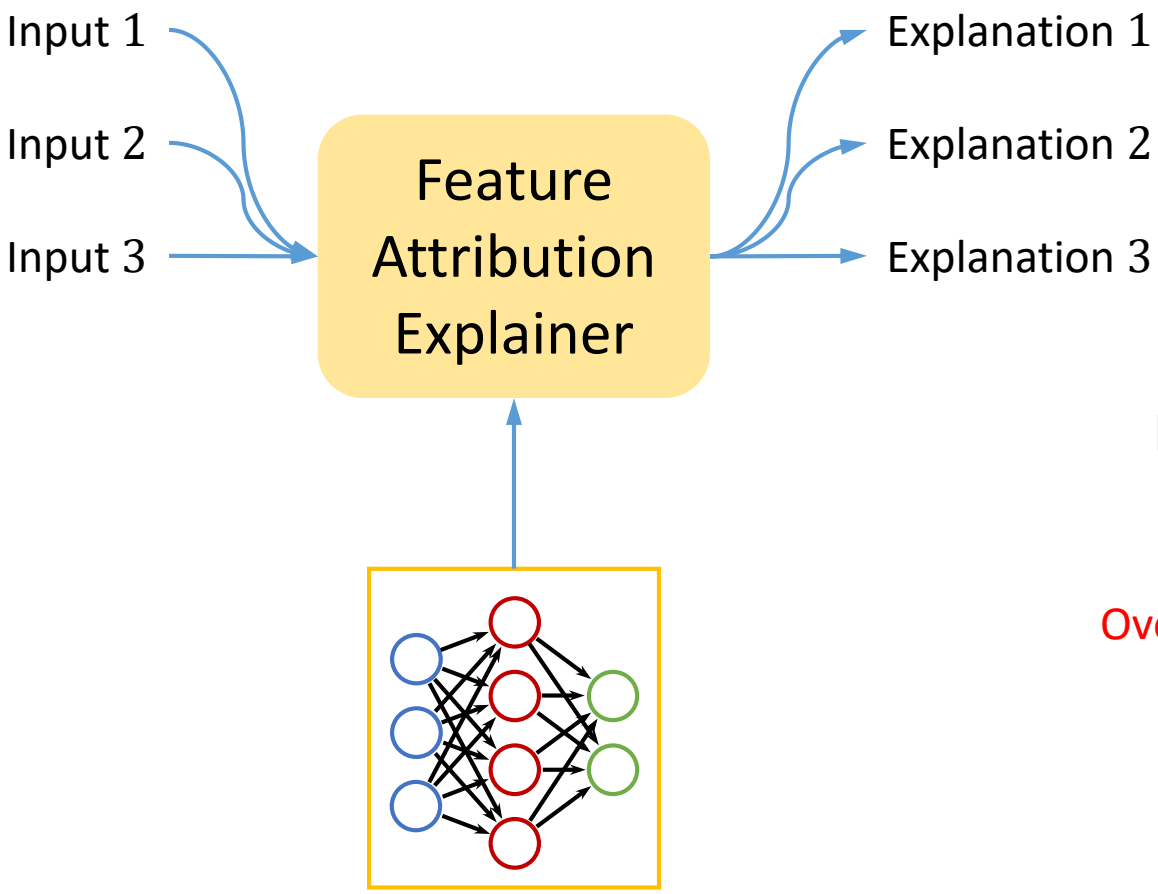

Over-generalization

Explanation Summary (ExSum)

Slides and Resources

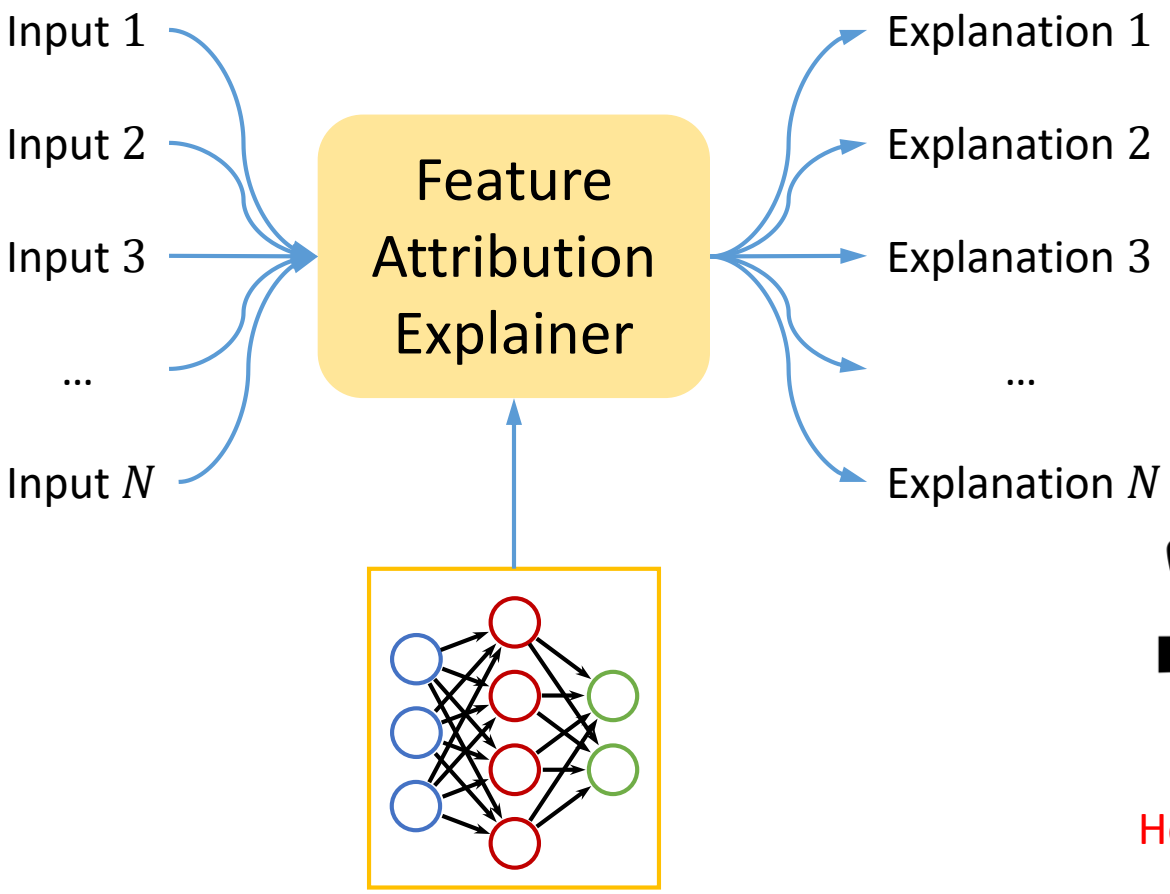

Heavy cognitive workload

Explanation Summary (ExSum)

Slides and Resources

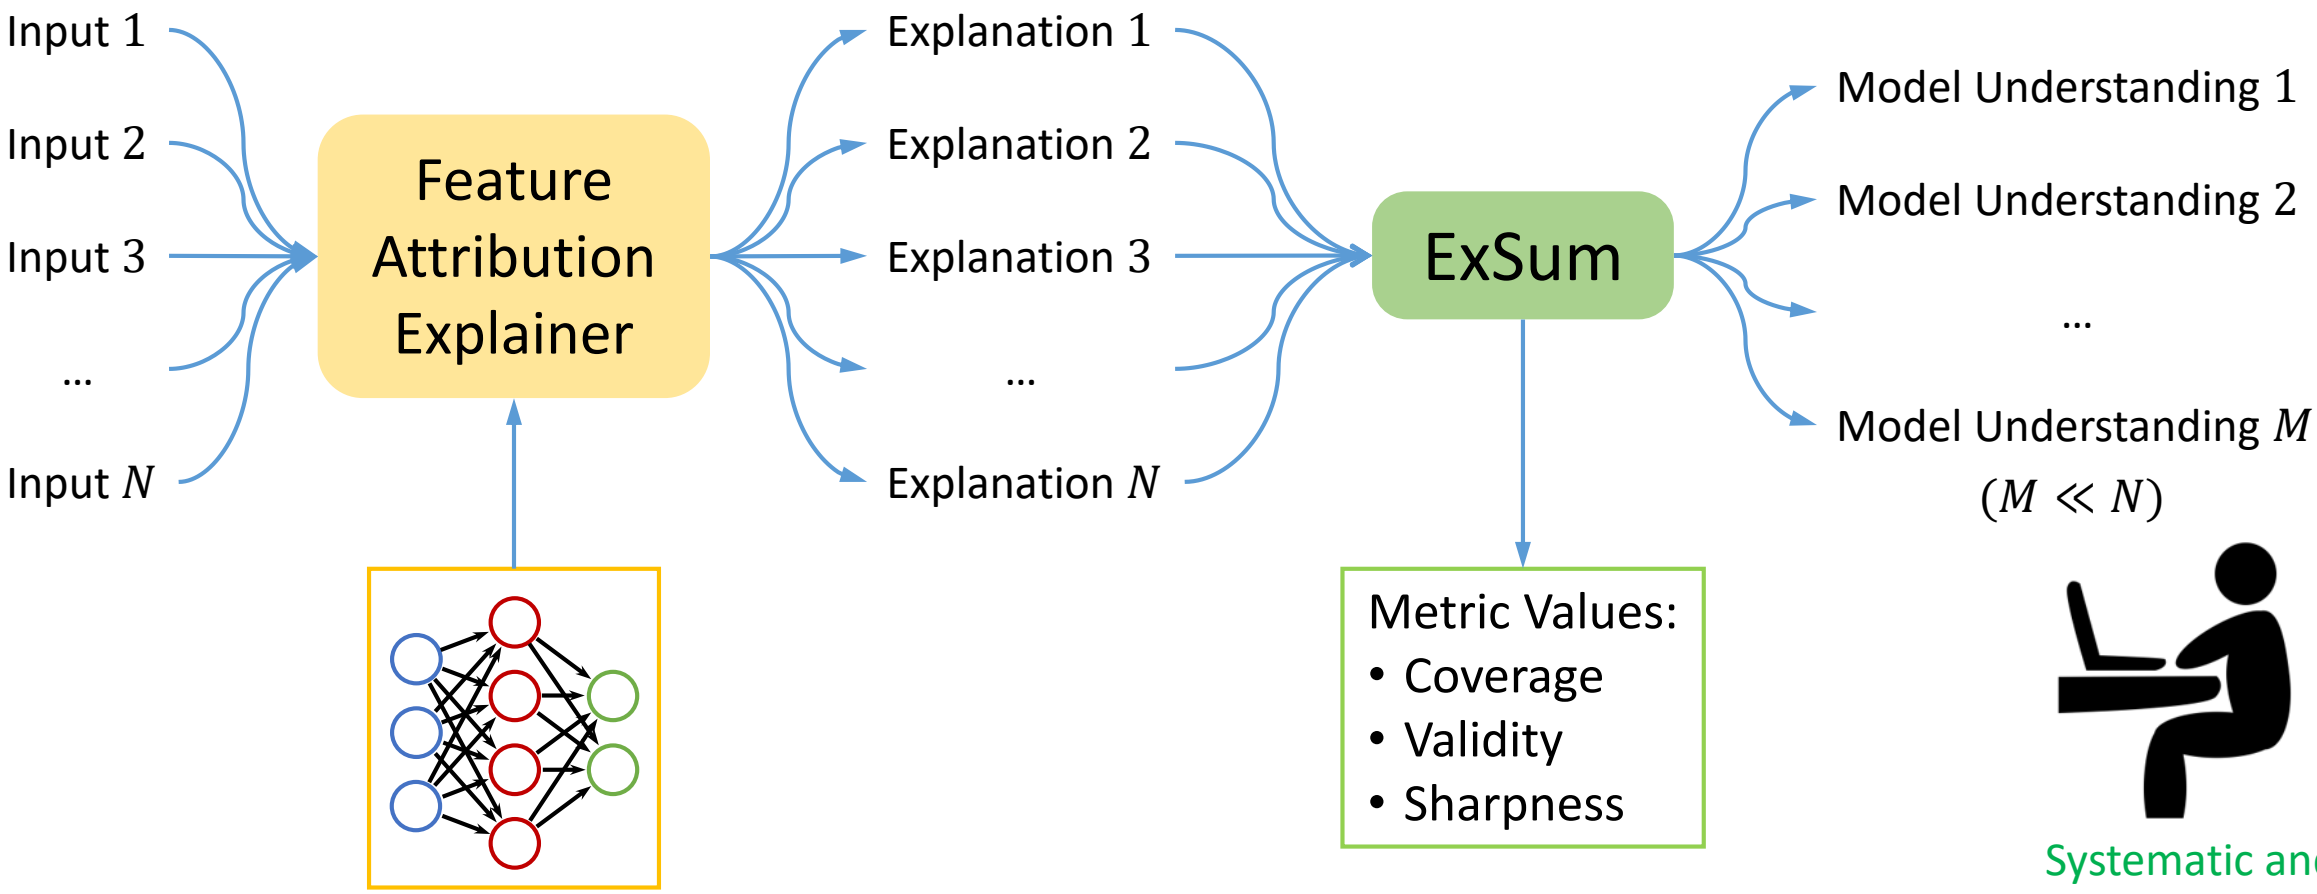

GUI for Developing ExSum Rules

Slides and Resources

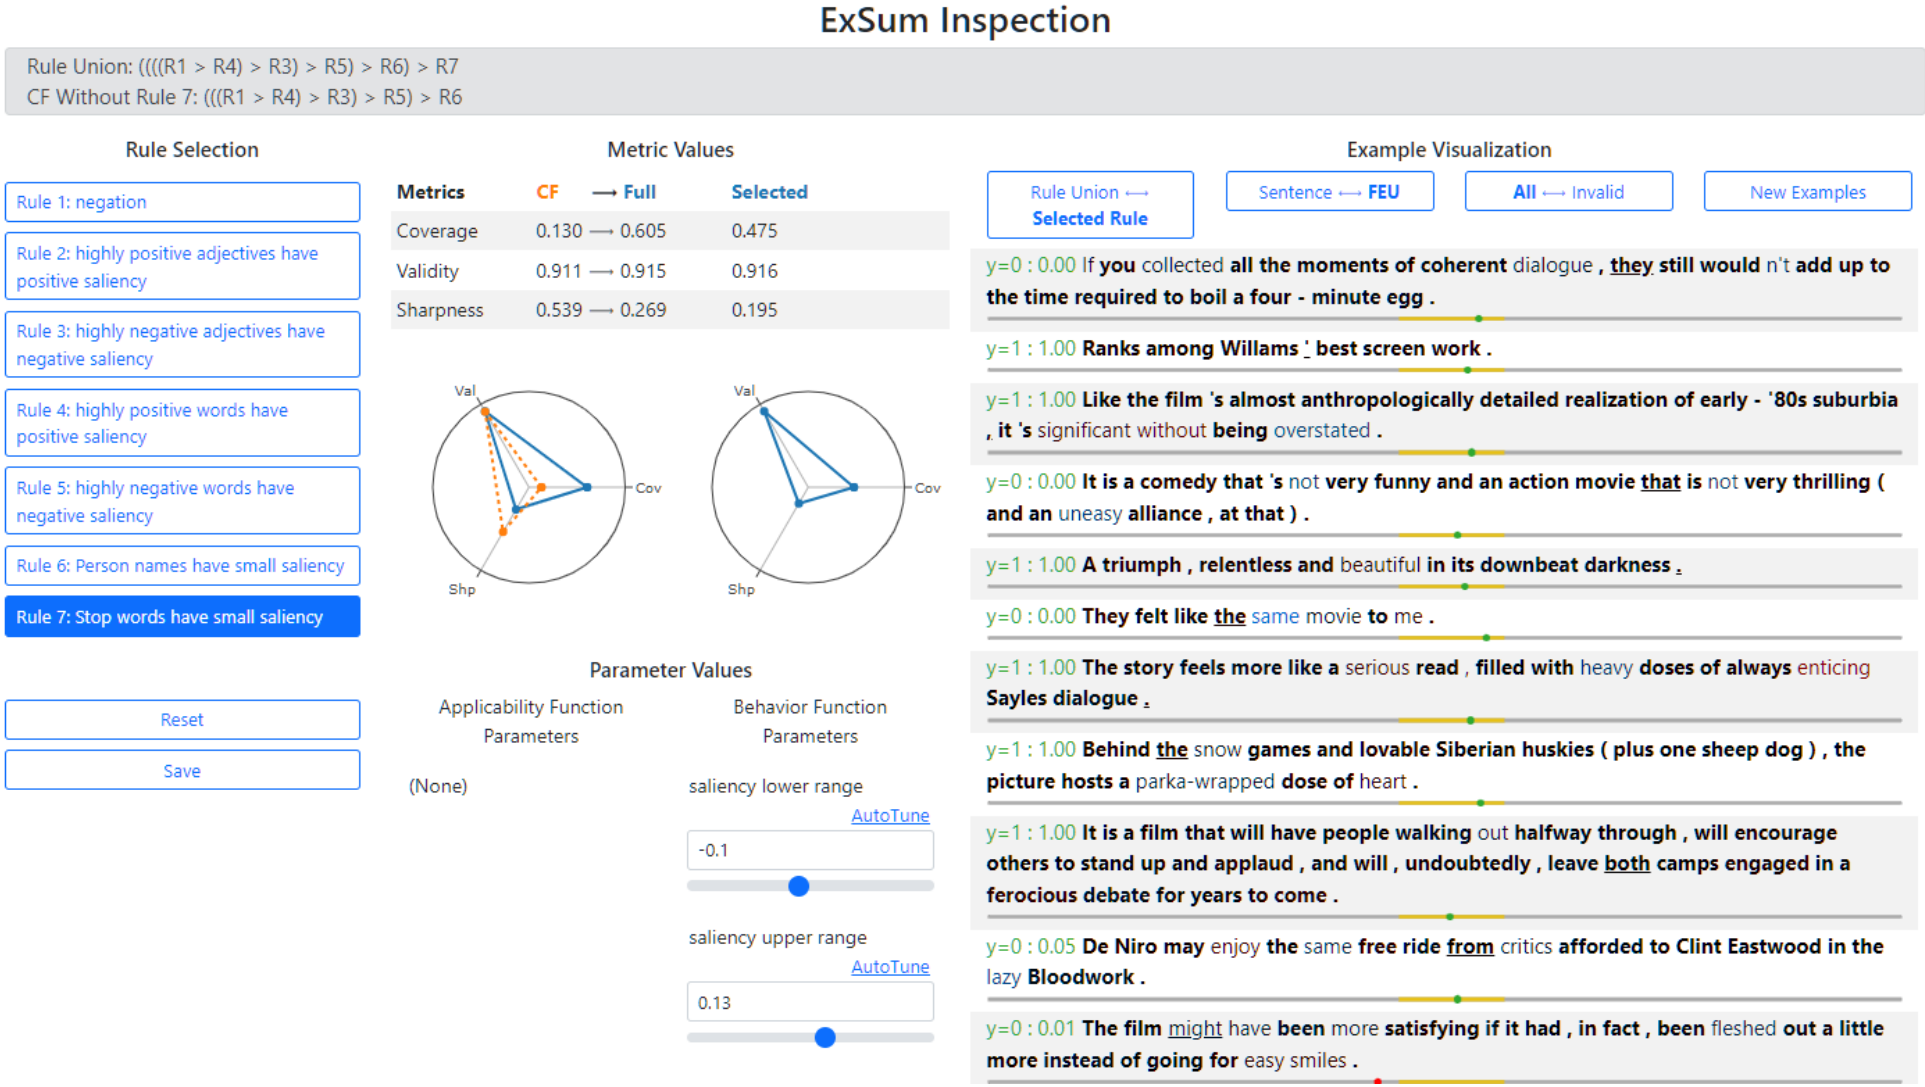

GUI for Developing ExSum Rules

Slides and Resources

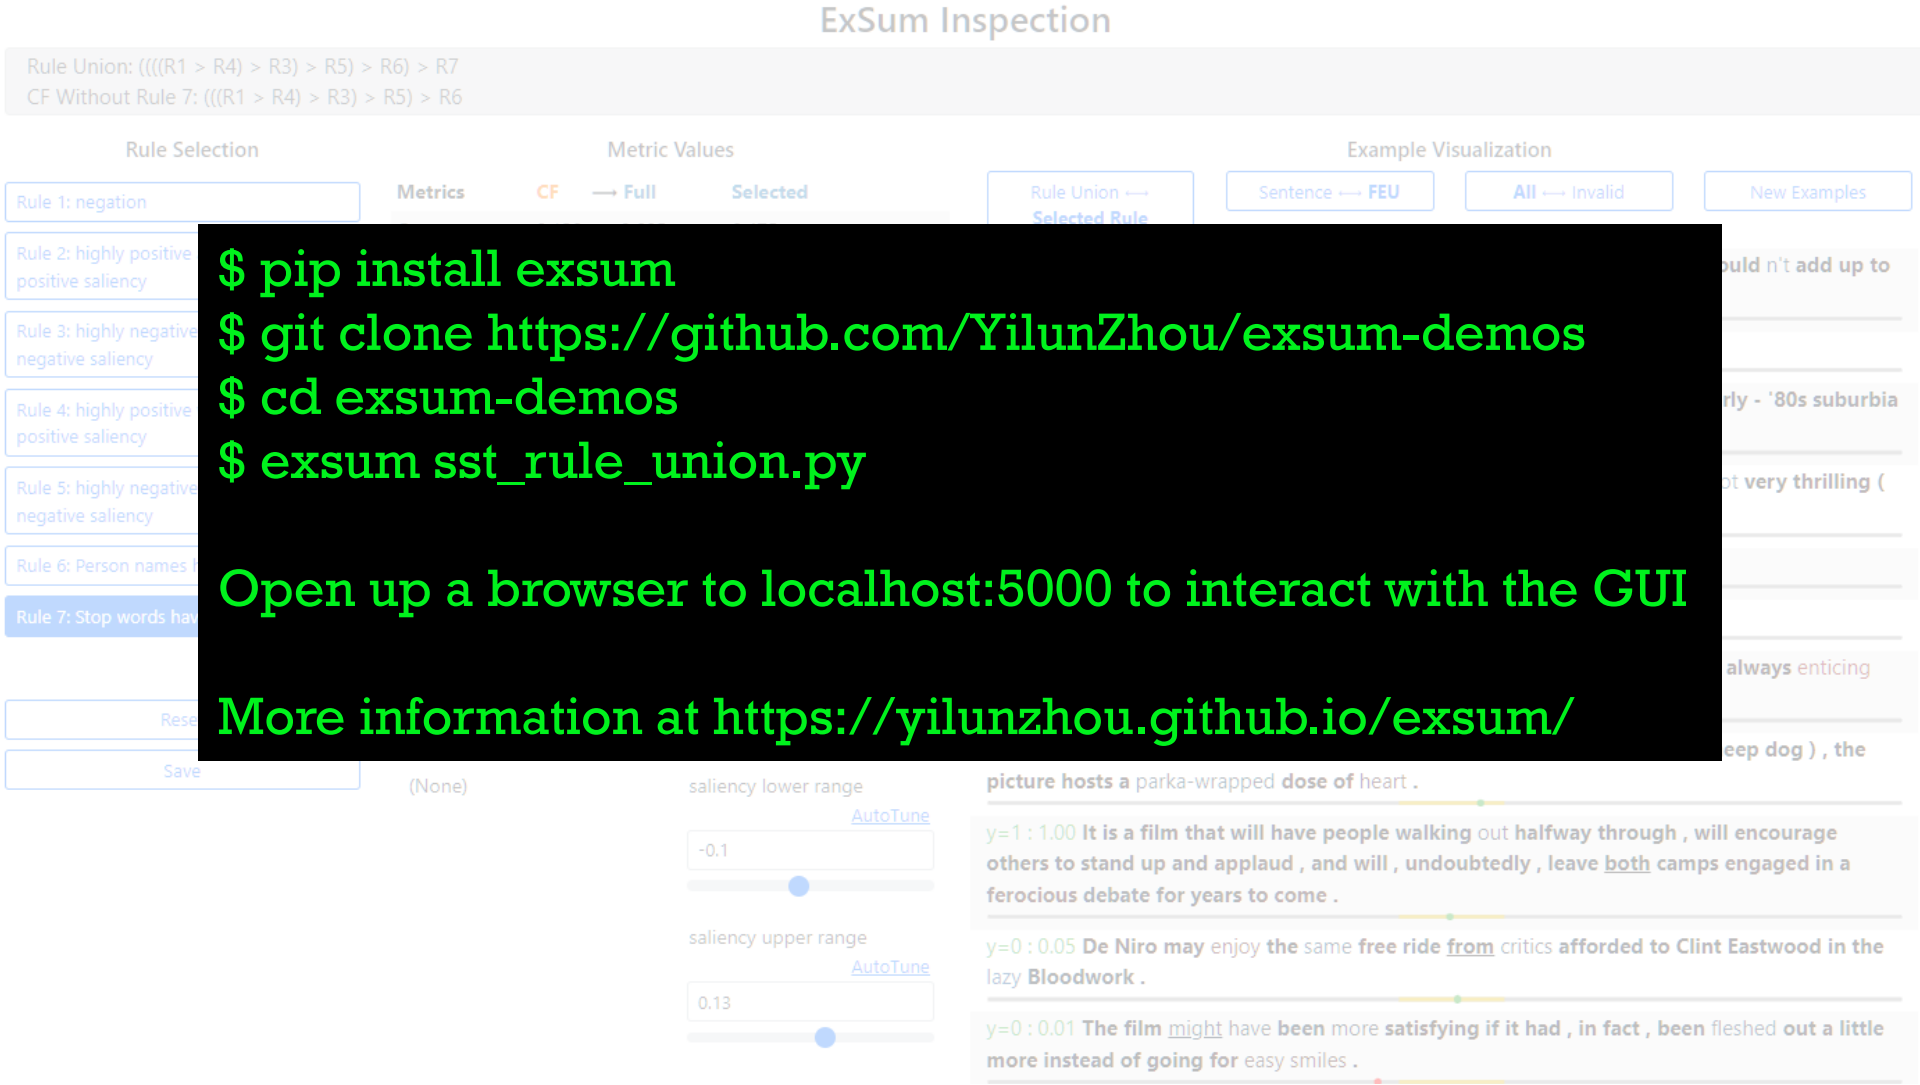


```
$ pip install exsum
$ git clone https://github.com/YilunZhou/exsum-demos
$ cd exsum-demos
$ exsum sst_rule_union.py

Open up a browser to localhost:5000 to interact with the GUI

More information at https://yilunzhou.github.io/exsum/
```

The Many Faces of Understandability

Slides and Resources

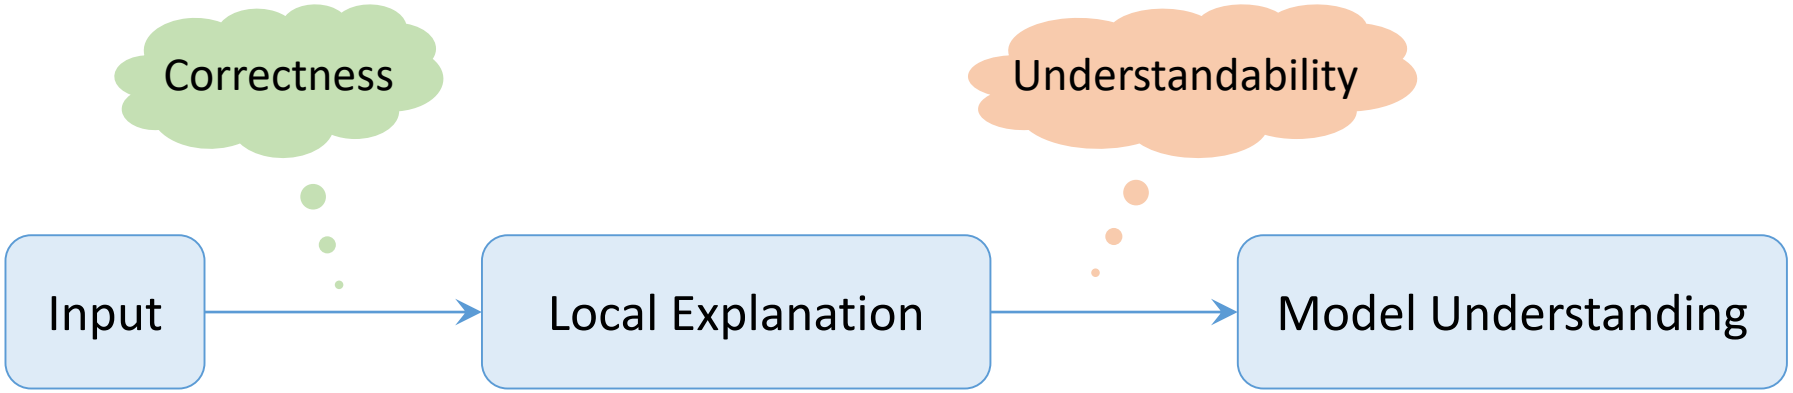

The Many Faces of Understandability

Slides and Resources

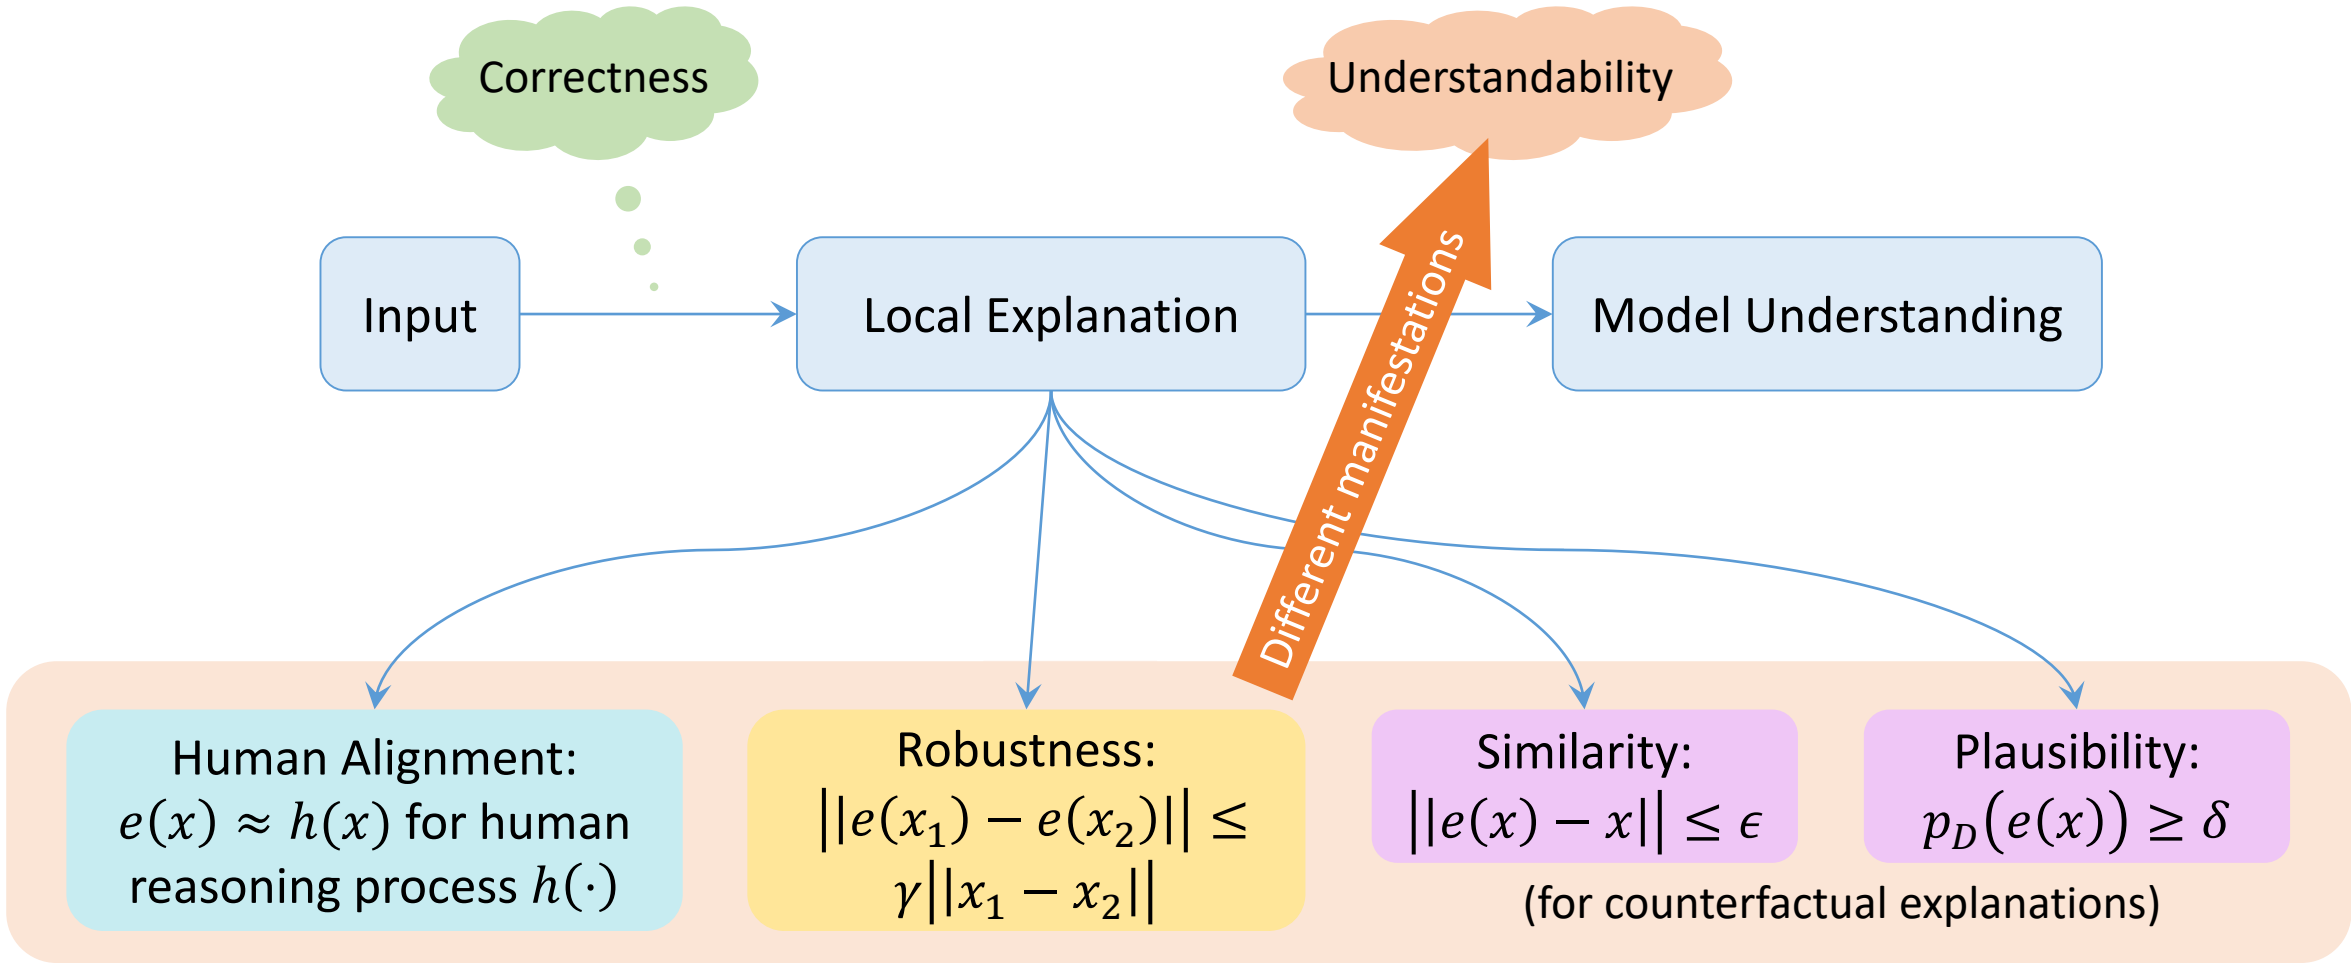

Definition vs. Evaluation

Slides and Resources

Definition

Gradient

$$D_g(x) = \nabla_x f(x)$$

Evaluation

Comprehensiveness

$$E_\kappa(x, e) = \frac{1}{L+1} \sum_{l=0}^{L} f(x) - f\left(\bar{x}_e^{(l)}\right)$$

Definition vs. Evaluation

Slides and Resources

Definition

Gradient

$$D_g(x) = \nabla_x f(x)$$

Evaluation

Comprehensiveness

$$E_\kappa(x, e) = \frac{1}{L+1} \sum_{l=0}^{L} f(x) - f\left(\bar{x}_e^{(l)}\right)$$

Gradient

$$E_g(x, e) = -\left|\left|\nabla_x f(x) - e\right|\right|$$

Definition vs. Evaluation

Slides and Resources

Definition

Gradient

$$D_g(x) = \nabla_x f(x)$$

Comprehensiveness

$$D_\kappa(x) = \underset{e}{\text{argmax}} \frac{1}{L+1} \sum_{l=0}^{L} f(x) - f\left(\bar{x}_e^{(l)}\right)$$

Evaluation

Comprehensiveness

$$E_\kappa(x, e) = \frac{1}{L+1} \sum_{l=0}^{L} f(x) - f\left(\bar{x}_e^{(l)}\right)$$

Gradient

$$E_g(x, e) = -\left|\left|\nabla_x f(x) - e\right|\right|$$

Definition-Evaluation Duality

Slides and Resources

Definition

Duality

Evaluation

Gradient

$$D_g(x) = \nabla_x f(x)$$

Comprehensiveness

$$D_\kappa(x) = \underset{e}{\text{argmax}} \frac{1}{L+1} \sum_{l=0}^{L} f(x) - f\left(\bar{x}_e^{(l)}\right)$$

Comprehensiveness

$$E_\kappa(x, e) = \frac{1}{L+1} \sum_{l=0}^{L} f(x) - f\left(\bar{x}_e^{(l)}\right)$$

Gradient

$$E_g(x, e) = -\left|\left|\nabla_x f(x) - e\right|\right|$$

Definition-Evaluation Duality

Slides and Resources

Definition

Duality

Evaluation

Gradient

$$D_g(x) = \nabla_x f(x)$$

Comprehensiveness

$$D_\kappa(x) = \underset{e}{\text{argmax}} \frac{1}{L+1} \sum_{l=0}^{L} f(x) - f\left(\bar{x}_e^{(l)}\right)$$

Comprehensiveness

$$E_\kappa(x, e) = \frac{1}{L+1} \sum_{l=0}^{L} f(x) - f\left(\bar{x}_e^{(l)}\right)$$

Gradient

$$E_g(x, e) = -\left|\left|\nabla_x f(x) - e\right|\right|$$

Definition-Evaluation Duality

Slides and Resources

Definition ⟷ Duality ⟷ Evaluation

Gradient

$$D_g(x) = \nabla_x f(x)$$

$D_g \gg_{\text{"ease"}} D_\kappa$

Comprehensiveness

$$D_\kappa(x) = \underset{e}{\text{argmax}} \frac{1}{L+1} \sum_{l=0}^{L} f(x) - f\left(\bar{x}_e^{(l)}\right)$$

Comprehensiveness

$$E_\kappa(x, e) = \frac{1}{L+1} \sum_{l=0}^{L} f(x) - f\left(\bar{x}_e^{(l)}\right)$$

Gradient

$$E_g(x, e) = -\left|\left|\nabla_x f(x) - e\right|\right|$$

Definition-Evaluation Duality

Slides and Resources

Definition ⟷ Duality ⟷ Evaluation

Gradient

$$D_g(x) = \nabla_x f(x)$$

$D_g \gg_{\text{"ease"}} D_\kappa$?

Comprehensiveness

$$D_\kappa(x) = \underset{e}{\text{argmax}} \frac{1}{L+1} \sum_{l=0}^{L} f(x) - f\left(\bar{x}_e^{(l)}\right)$$

Comprehensiveness

$$E_\kappa(x, e) = \frac{1}{L+1} \sum_{l=0}^{L} f(x) - f\left(\bar{x}_e^{(l)}\right)$$

Gradient

$$E_g(x, e) = -\left|\left|\nabla_x f(x) - e\right|\right|$$

Beam Search Results

Slides and Resources

$$\kappa(x, e) = \frac{1}{L+1}\sum_{l=0}^{L} f(x) - f\left(\bar{x}_e^{(l)}\right)$$

$$e^* = \underset{e}{\operatorname{argmax}}\, \kappa(x, e)$$

Algorithm 1: Beam search for finding e^*.

Input: beam size B, metric m, sentence x of length L;
Let $e^{(0)}$ be an empty length-L explanation;
beams $\leftarrow \{e^{(0)}\}$;
for $l = 1, ..., L$ **do**
  beams $\leftarrow \bigcup_{e \in \text{beams}} \text{ext}(e, L - l + 1)$;
  beams $\leftarrow$ choose_best(beams, B);
end
$e^* \leftarrow$ choose_best(beams, 1);
$e^* \leftarrow \text{shift}(e^*)$;
return e^*;

Beam Search Results

Slides and Resources

$$\kappa(x, e) = \frac{1}{L+1}\sum_{l=0}^{L} f(x) - f\left(\bar{x}_e^{(l)}\right)$$

$$e^* = \underset{e}{\operatorname{argmax}}\, \kappa(x, e)$$

Algorithm 1: Beam search for finding e^*.

Input: beam size B, metric m, sentence x of length L;
Let $e^{(0)}$ be an empty length-L explanation;
$\texttt{beams} \leftarrow \{e^{(0)}\}$;
for $l = 1, ..., L$ **do**
  $\texttt{beams} \leftarrow \bigcup_{e \in \texttt{beams}} \text{ext}(e, L - l + 1)$;
  $\texttt{beams} \leftarrow \text{choose\_best}(\texttt{beams}, B)$;
end
$e^* \leftarrow \text{choose\_best}(\texttt{beams}, 1)$;
$e^* \leftarrow \text{shift}(e^*)$;
return e^*;

A worthy tribute to a great humanitarian and her vibrant ‘ co-stars . ’

So stupid , so ill-conceived , so badly drawn , it created whole new levels of ugly .

Beam Search Results

Slides and Resources

$$\kappa(x, e) = \frac{1}{L+1}\sum_{l=0}^{L} f(x) - f\left(\bar{x}_e^{(l)}\right)$$

$$e^* = \underset{e}{\operatorname{argmax}}\, \kappa(x, e)$$

Algorithm 1: Beam search for finding e^*.

Input: beam size B, metric m, sentence x of length L;
Let $e^{(0)}$ be an empty length-L explanation;
beams $\leftarrow \{e^{(0)}\}$;
for $l = 1, ..., L$ **do**
  beams $\leftarrow \bigcup_{e \in \text{beams}} \text{ext}(e, L - l + 1)$;
  beams $\leftarrow$ choose_best(beams, B);
end
$e^* \leftarrow$ choose_best(beams, 1);
$e^* \leftarrow$ shift(e^*);
return e^*;

A worthy tribute to a great humanitarian and her vibrant ‘ co-stars . ’

So stupid , so ill-conceived , so badly drawn , it created whole new levels of ugly .

Explainer	Comp $\kappa \uparrow$
Grad	0.327
IntG	0.525
LIME	0.682
SHAP	0.612
Occl	0.509
E*	0.740
Random	0.218

Proxy Metrics for Explanation Quality

Slides and Resources

- Feature importance ⟺ model prediction change with feature removal

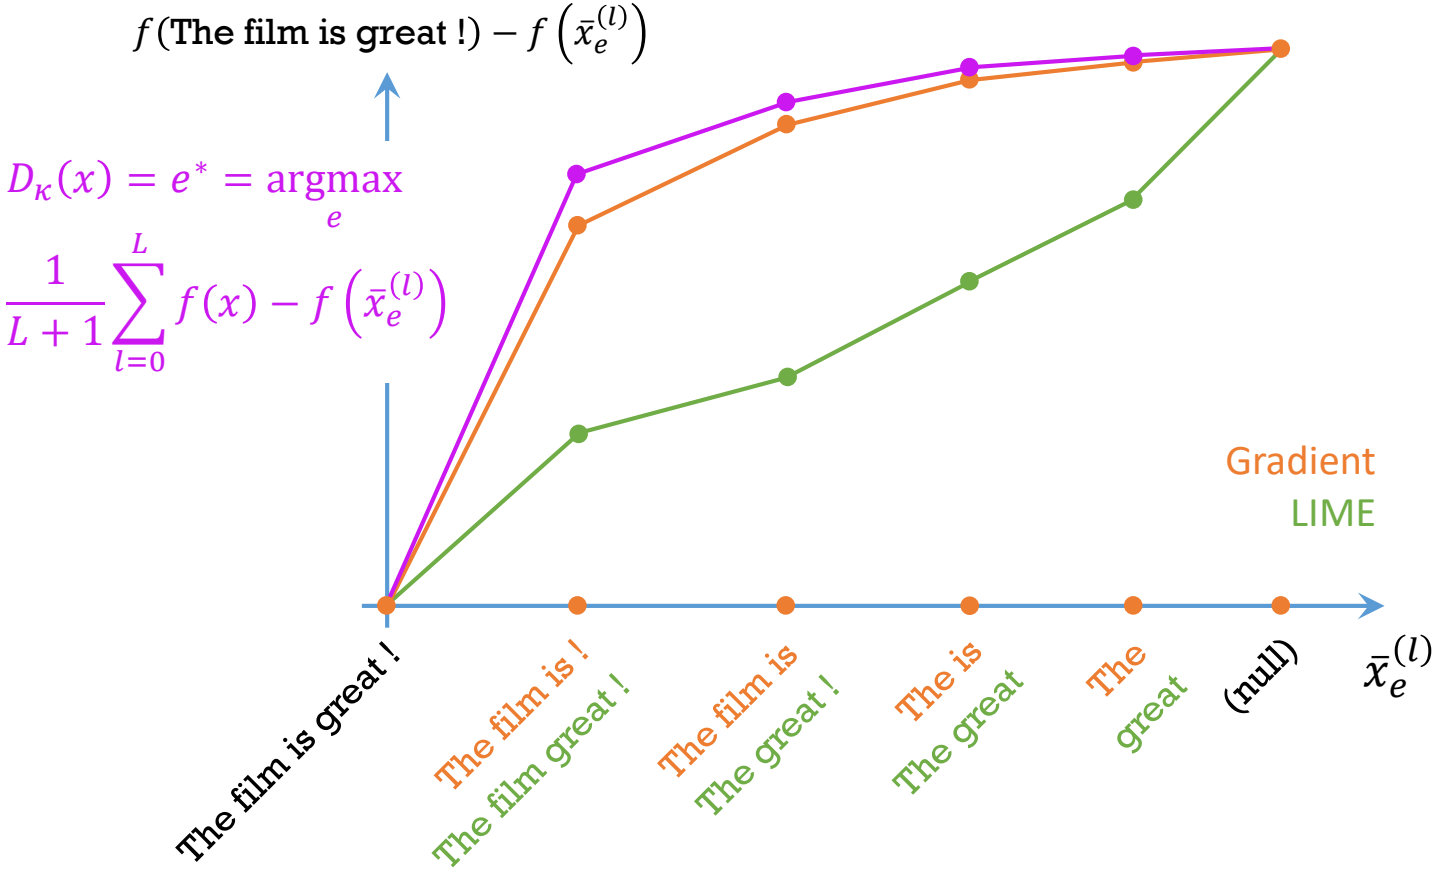

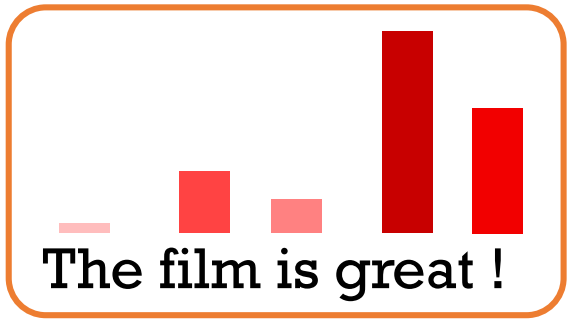

Gradient

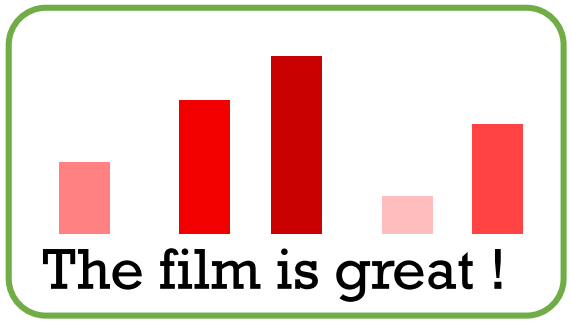

LIME

Evaluation on Other Aspects

Slides and Resources

- E* is competitive on other metrics

Explainer	$DF_{MIT}\uparrow$	$DF_{Frac}\downarrow$	$Rank_{Del}\uparrow$
Grad	10.5%	54.5%	0.162
IntG	16.9%	39.6%	0.369
LIME	25.5%	28.1%	0.527
SHAP	23.0%	36.1%	0.369
Occl	26.4%	40.6%	1.000
E*	25.0%	25.2%	0.438
Random	3.4%	72.3%	0.004

Evaluation on Other Aspects

Slides and Resources

- E* is competitive on other metrics
- E* is robust to perturbations

Explainer	$DF_{MIT}\uparrow$	$DF_{Frac}\downarrow$	$Rank_{Del}\uparrow$
Grad	10.5%	54.5%	0.162
IntG	16.9%	39.6%	0.369
LIME	25.5%	28.1%	0.527
SHAP	23.0%	36.1%	0.369
Occl	26.4%	40.6%	1.000
E*	25.0%	25.2%	0.438
Random	3.4%	72.3%	0.004

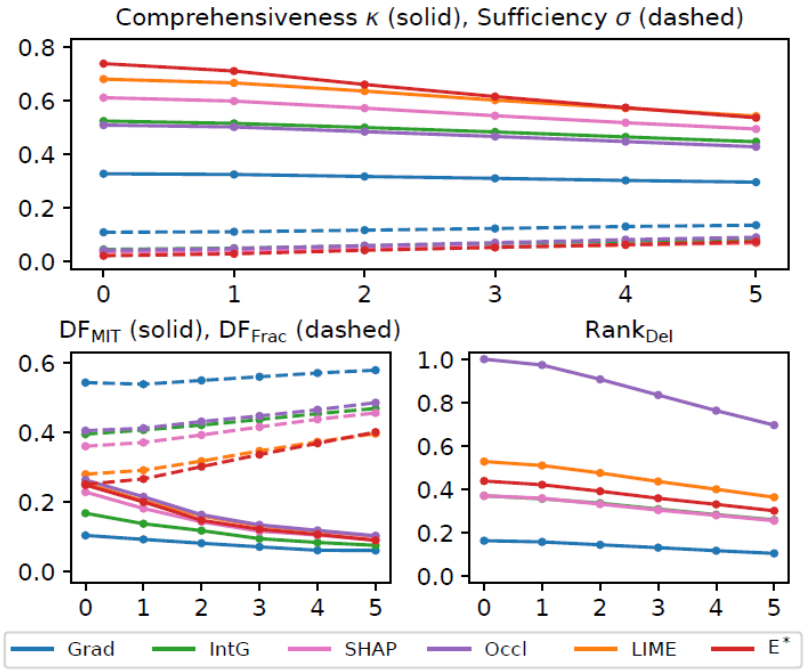

A New Paradigm of Developing Explainers?

Slides and Resources

A New Paradigm of Developing Explainers?

Slides and Resources

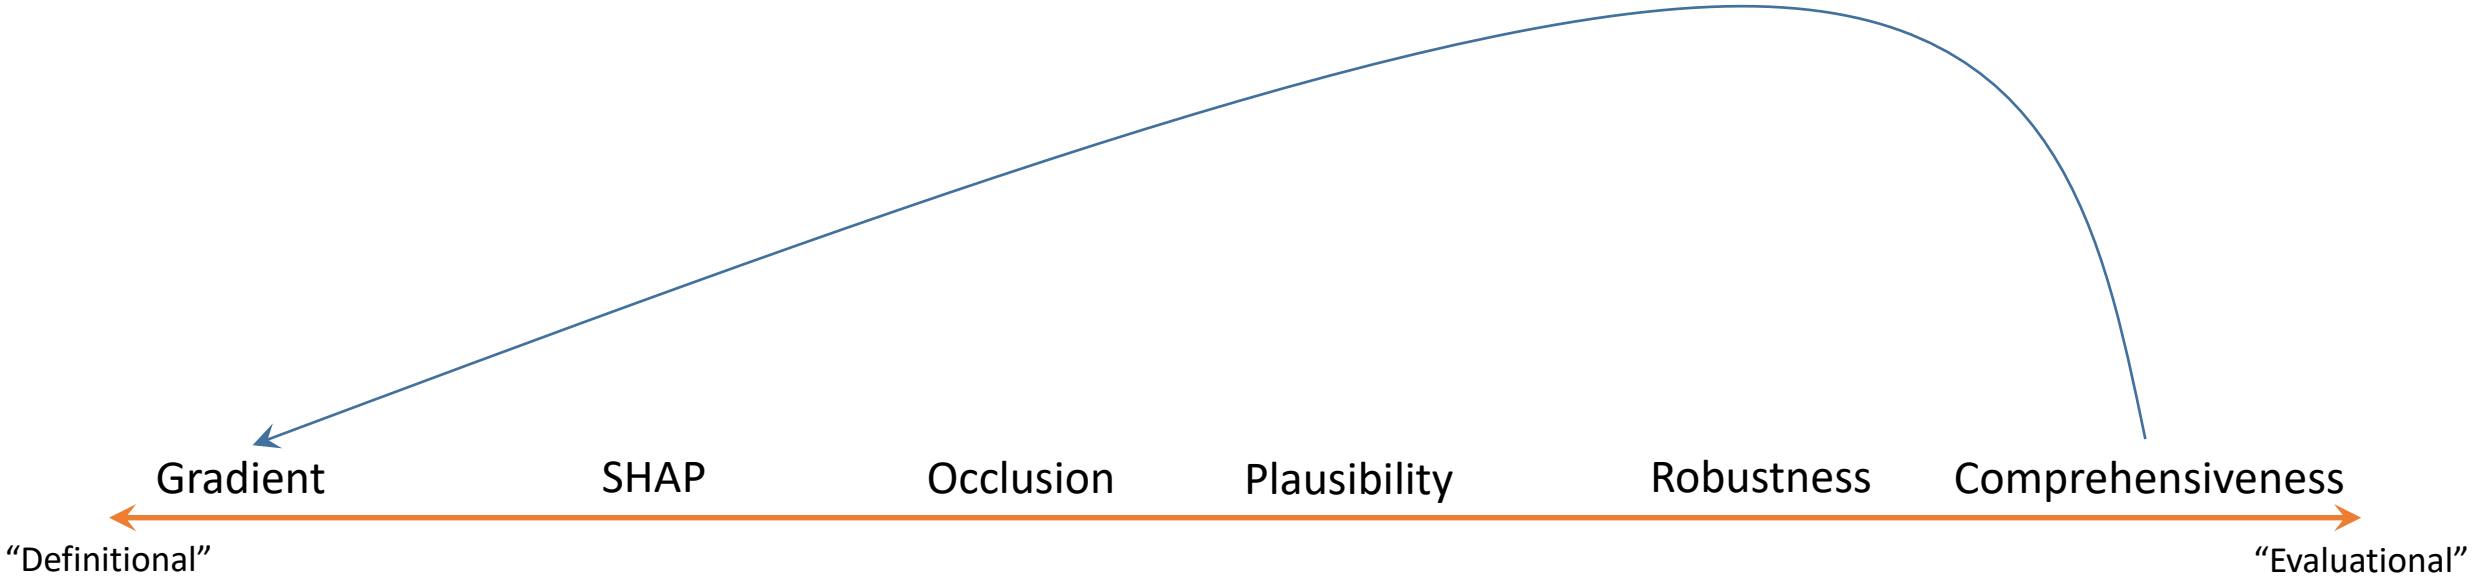

A New Paradigm of Developing Explainers?

Slides and Resources

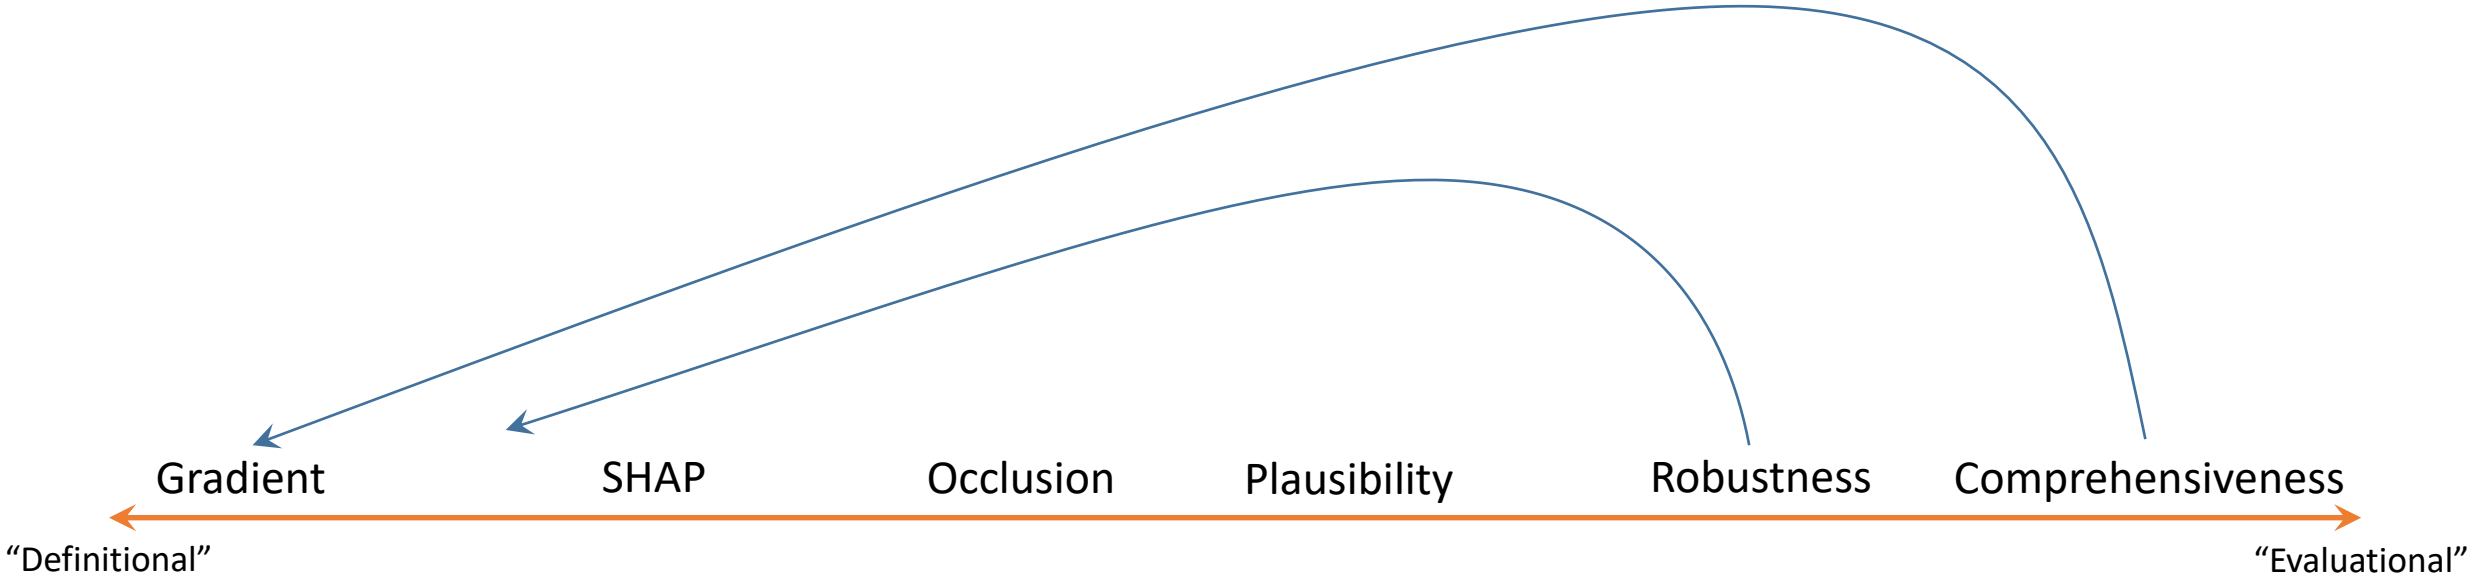

A New Paradigm of Developing Explainers?

Slides and Resources

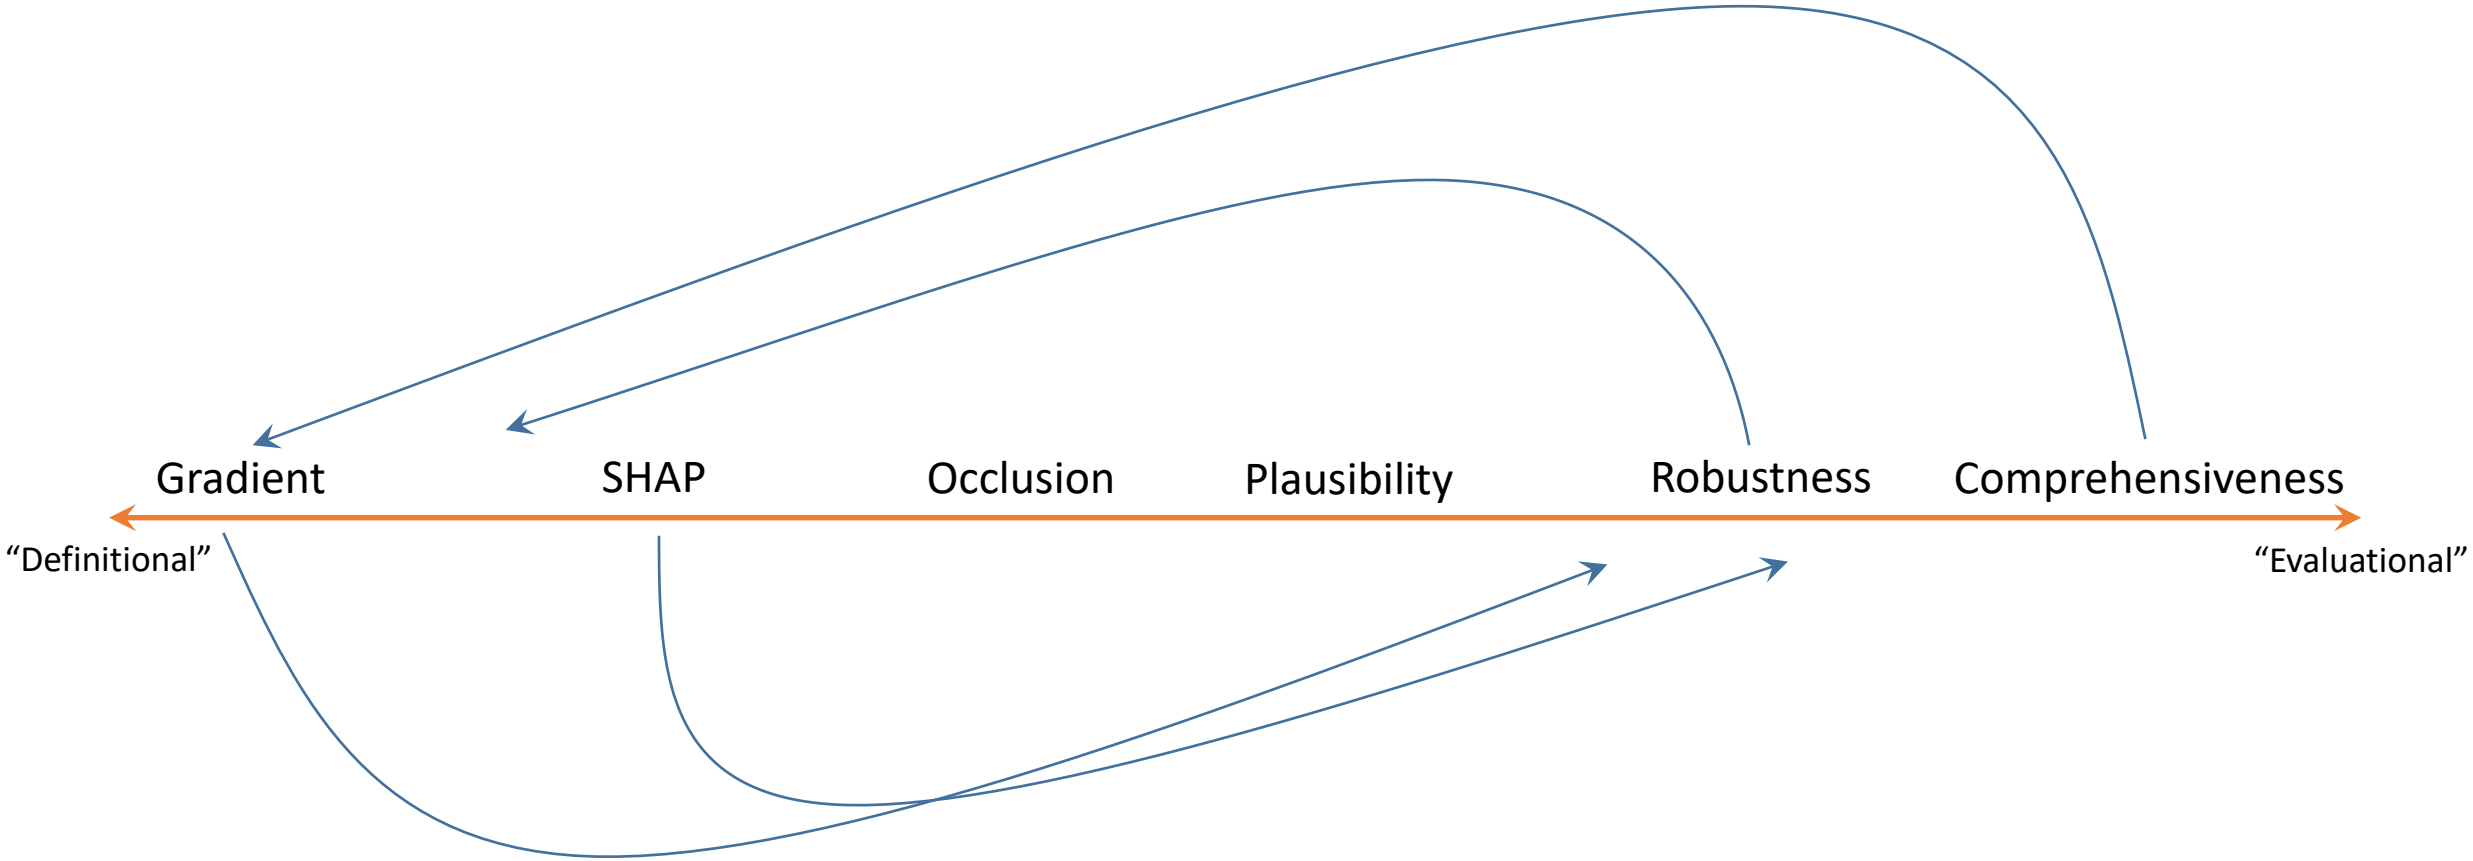

Solvability-Based Explainer

Slides and Resources

pip install solvex

https://yilunzhou.github.io/solvability/

Summary | NLP - Word | NLP - Sentence | CV - Grid Superpixel | CV - Custom Superpixel | Tabular

Read this tutorial as a Jupyter notebook here.

In this demo, we compute word-level explanations for the Huggingface `textattack/roberta-base-SST-2` model, also the setup presented in the paper.

We first load required packages and the RoBERTa model. Two classes are needed to compute the explanations. `BeamSearchExplainer` implements the beam search algorithm, and `*Masker` implements the feature masking. In this demo, we use `TextWordMasker` since we need to mask out individual words from a text input. The other demos showcase other `*Masker`s.

```
from solvex import BeamSearchExplainer, TextWordMasker
import torch
from transformers import AutoTokenizer, AutoModelForSequenceClassification
```

Solvability-Based Explainer

Slides and Resources

pip install solvex

Explained label: 1

Function value for label 1: 1.000

Contrary to other reviews, I have zero complaints about the service or the prices. I have been getting tire service here for the past 5 years now, and compared to my experience with places like Pep Boys, these guys are experienced and know what they're doing. Also, this is one place that I do not feel like I am being taken advantage of, just because of my gender. Other auto mechanics have been notorious for capitalizing on my ignorance of cars, and have sucked my bank account dry. But here, my service and road coverage has all been well explained - and let up to me to decide. And they just renovated the waiting room. It looks a lot better than it did in previous years.

Solvability-Based Explainer

Slides and Resources

pip install solvex

Explained label: 1

Function value for label 1: 1.000

Contrary to other reviews, I have zero complaints about the service or the prices. I have been getting tire service here for the past 5 years now, and compared to my experience with places like Pep Boys, these guys are experienced and know what they're doing. Also, this is one place that I do not feel like I am being taken advantage of, just because of my gender. Other auto mechanics have been notorious for capitalizing on my ignorance of cars, and have sucked my bank account dry. But here, my service and road coverage has all been well explained - and let up to me to decide. And they just renovated the waiting room. It looks a lot better than it did in previous years.

Explained label: 232. Function value: 0.159

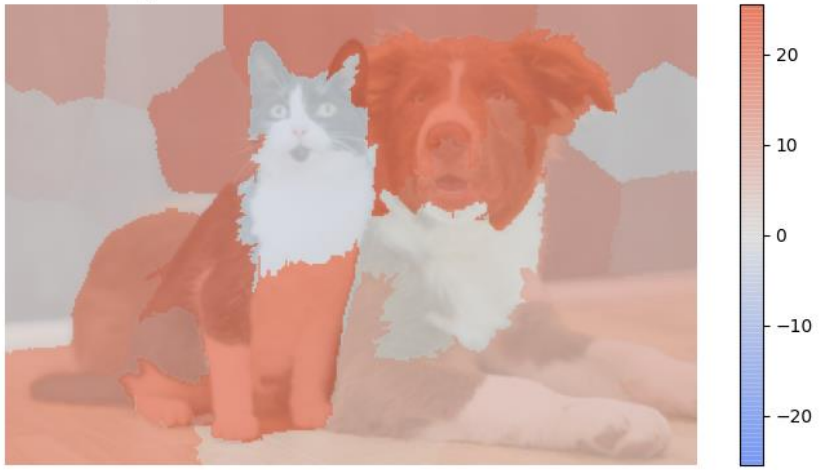

Solvability-Based Explainer

Slides and Resources

pip install solvex

Explained label: 1

Function value for label 1: 1.000

Contrary to other reviews, I have zero complaints about the service or the prices. I have been getting tire service here for the past 5 years now, and compared to my experience with places like Pep Boys, these guys are experienced and know what they're doing. Also, this is one place that I do not feel like I am being taken advantage of, just because of my gender. Other auto mechanics have been notorious for capitalizing on my ignorance of cars, and have sucked my bank account dry. But here, my service and road coverage has all been well explained - and let up to me to decide. And they just renovated the waiting room. It looks a lot better than it did in previous years.

Explained label: 232. Function value: 0.159

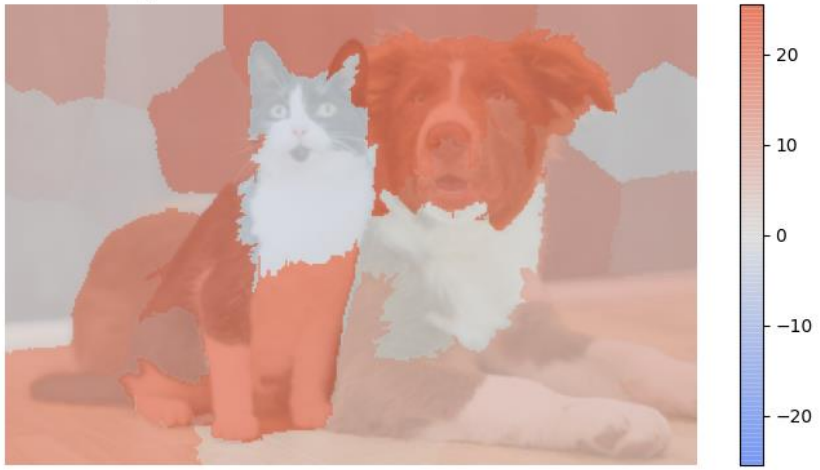

Explained label: 0. Function value: 0.980

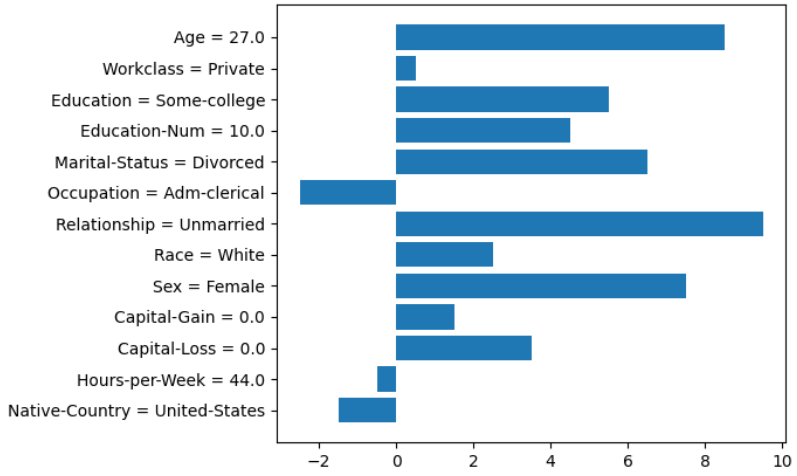

References

Slides and Resources

- Definitions
  - Simonyan et al. Deep Inside Convolutional Networks: Visualising Image Classification Models and Saliency Maps. arXiv 2014
  - Zeiler and Fergus. Visualizing and Understanding Convolutional Networks. ECCV 2014
  - Ribeiro et al. “Why Should I Trust You?” Explaining the Predictions of Any Classifier. KDD 2016
  - Smilkov et al. SmoothGrad: Removing Noise by Adding Noise. arXiv 2017
  - Sundararajan et al. Axiomatic Attribution for Deep Networks. ICML 2017
  - Lundberg et al. A Unified Approach to Interpreting Model Predictions. NIPS 2017
- Evaluations
  - Samek et al. Evaluating the Visualization of What a Deep Neural Network Has Learned. T-NNLS 2016
  - Petsiuk et al. RISE: Randomized Input Sampling for Explanation of Black-box Models. BMVC 2018
  - Ghorbani et al. Interpretation of Neural Networks is Fragile. AAAI 2019
  - DeYoung et al. ERASER: A Benchmark to Evaluate Rationalized NLP Models. ACL 2020
  - Ross et al. Explaining NLP Models via Minimal Contrastive Editing (MiCE). ACL 2021 (Findings)
  - Zhou et al. Do Feature Attribution Methods Correctly Attribute Features? AAAI 2022
  - Zhou et al. ExSum: From Local Explanations to Model Understanding. NAACL 2022
- Duality
  - Zhou and Shah. The Solvability of Interpretability Evaluation Metrics. EACL 2023 (Findings)